

TWIN-64

Architecture Document

Helmut Fieres
August 20, 2025

Contents

Introduction	1
System Organization	3
A RegisterMemory Architecture	3
Processor	3
Virtual Memory	4
Security and Protection Support	4
Privilege Changes	4
Debugging Support	5
Input/Output	5
Interrupt System	5
Instruction Set	7
Computational Instructions	7
Immediate Instructions	7
Memory Reference Instructions	8
Control Flow Instructions	8
System Control Instructions	8
Processing Resources	9
General Registers	9
Control Registers	9
Processor State Register	12
Data Types	13
Byte Ordering	14
Translation Lookaside Buffer	14
Caches	15
System Memory	17
Virtual Address Space	17
Physical Address Space	18
Addressing and Access Control	19
Address Translation	19
Access Control	19
Access Protection	20
Control Flow	23
Unconditional Branches	23
Conditional Branches	23
Linkage	23
Branching and Segments	24
Privilege Level Changes	24

CONTENTS

Traps Associated with Branches	24
Interruptions	25
Interrupt Handling	25
Interrupt Vector Table	26
Instruction Set Overview	27
Instruction Formats	27
Arithmetic Instructions	28
Bit Operations Instructions	28
Immediate Instructions	28
Memory Reference Instructions	29
Control Flow Instructions	29
System Instructions	30
Instruction Set Descriptions	31
ABR - Add and branch	32
ADD - Integer Addition	33
ADDIL Add immediate left	34
AND Boolean Operation	35
B branch unconditional	36
BB branch on bit	37
BR branch register	38
BV branch vectored	39
CBR compare and branch	40
CMP Integer Compare	41
DEP Deposit Bit Field	42
DIAG -	43
DSR Double Shift Right	44
EXTR Extract Bit Field	45
FCA - flush from cache	46
ITLB insert TLB entry	47
LD Load register	48
LDI Load immediate	49
LDO load offset	50
LDR Load register reserved	51
LPA - load physical address	52
MBR move and branch	53
MFIA - Move from instruction address register	54
MFCR Move from Control register	55
MFCR Move to Control register	56
NOP - No operation	57
OR Boolean Operation	58
PCA - purge from cache	59
PRB - probe virtual address	60
PTLB - purge from TLB	61
RFI - Return from interrupt	62
RSM - Reset system mask	63
SHLxA Shift left and add	64
SHRxA Shift right and add	65
SSM - Set system mask	66
ST Store register	67

ST Store register conditional	68
SUB - Integer Subtraction	69
DIAG -	70
XOR Boolean Operation	71
Input/Output	73
Concepts	73
I/O Address Space	73
Hard physical address range	74
Broadcast address range	75
Soft physical address range	75
I/O Elements	75
Appendix - Instruction Commentary	77
xxx	77
Appendix - Instruction Notation Routines	79
Appendix - TLB and Caches	81
xxx	81

Introduction

Designers of CPUs in the seventies and early eighties almost all used a microcoded approach with hundreds of instructions. RISC was just beginning to emerge. Registers were typically not fully general-purpose but often had special functions or meanings, and many instructions were rather complicated, though designers believed them useful. Compiler writers, however, often used only a small subset, ignoring these complex instructions because they were too specialized. In the late eighties and nineties the shift moved to RISC-based designs, with large sets of registers, fixed instruction word lengths, large virtual memory address ranges, and instructions that were in general much more pipeline-friendly. Today, processors are highly complex, with superscalar deep pipelines and multilevel cache hierarchies—just a few of the hallmarks of modern CPU design. CPU designs have become so large and complex that no single person can completely understand them. But what if there were a simple design that one person could fully understand?

Welcome to Twin-64. Twin-64 is a simple 64-bit CPU with a register-memory model and segmented virtual memory. The design is heavily influenced by Hewlett-Packard's PA-RISC architecture, which was initially a 32-bit RISC-style load/store machine. Almost all of the key architectural features come from there. However, other processors such as the Data General MV8000, the DEC Alpha, and the IBM PowerPC were also influential or at least investigated for their design choices. In addition, the Blitz-64 architecture was studied. Blitz-64, a design developed at Portland State University, offered an interesting approach in that it only supported 64-bit signed arithmetic, avoiding common errors from mixing signed and unsigned arithmetic.

While it is not a design goal to build an industry-strength, competitive CPU, the CPU should nevertheless be useful and largely follow established design practices. For example, instructions should not be invented just because they seem useful, but rather designed with the assembler and compilers in mind. Instruction set design is CPU design. Register-memory architectures were typically microcoded, complex-instruction machines. In contrast, Twin-64 instructions will be hardwired and are in general pipeline-friendly, avoiding data and control hazards and stalls where possible.

Where does the name Twin-64 come from? Obviously, the 64 indicates that the CPU is a 64-bit CPU. The Twin part will become clearer when implementations of the CPU are described in later chapters. One implementation represents a dual-issue instruction design with two ALUs for computation. Nevertheless, a very basic implementation could just offer single-issue, single-ALU execution as well.

This document describes the overall architecture, instruction set, and runtime model for Twin-64. It is structured into several parts. The first part gives an overview of the system organization and architecture. The main part of the document presents the instruction set. Memory reference instructions, immediate instructions, branch instructions, computational instructions, and system control instructions are described in detail. These chapters are the authoritative reference for the instruction set. The next part presents the major components such as the processor itself, the I/O module architecture, and finally the runtime model and calling conventions. An appendix provides additional design details and complementary notes.

System Organization

This chapter introduces Twin-64 by highlighting the major components and design decisions taken.

A RegisterMemory Architecture

Modern RISC CPUs are typically load/store designs and feature a large number of general-purpose registers. Operations take place between registers. Twin-64 follows a slightly different route. It has fewer general registers, and in addition to registerregister computation, a register-memory model is also supported. However, there are no instructions that read from and write to memory in a single instruction cycle, since this would complicate the design considerably.

Twin-64 implements a **register-memory architecture** offering only a few addressing modes for the operand, combined with a fixed instruction word length. For most instructions one operand is a register, while the other is a flexible operand that can be an immediate, another register, or a memory address from which data is obtained or stored. The result is placed in the register, which is also the first operand. These types of machines are also called two-address machines.

In contrast to similar historical registermemory designs, Twin-64 has no operations that both read and write memory in the same instruction. For example, there is no increment memory instruction, since this would require two memory operations and is generally not pipeline-friendly. Memory data access is performed at the word, half-word, or byte level. Besides the implicit operand fetch in a computational instruction, there are dedicated load/store instructions. In addition to the registerimmediate and registermemory operand modes, one operand mode supports the three-operand model ($R_s = R_a \text{ OP } R_b$), specifying two source registers and a result register. This allows Twin-64 to support both registermemory and load/store operation styles.

Processor

The Twin-64 processor is only one element of a complete system. A system also includes memory, I/O modules, adapters, and interconnecting buses. The Central Processing Unit (CPU) includes a general register set and machine state registers. To support virtual memory addressing, a hardware translation lookaside buffer (TLB) is included to provide virtual-to-physical address translation. Instruction and data caches are optional, but for performance reasons are also key components of the processor. A processor will always emit virtual addresses that must be translated into physical addresses by the TLB. A Twin-64 system can consist of multiple processors.

Virtual Memory

Twin-64 offers a large global address range, organized in segments. A segment number and offset together form the global virtual address. Segments are also associated with access rights and protection identifiers. The CPU itself can run in user or privileged mode. Twin-64 always generates virtual addresses at the CPU level. The global virtual memory range is a 52-bit space organized into a 20-bit segment ID and a 32-bit byte offset. There are 2^{20} segments with 2^{32} bytes each.

Virtual addresses are translated to physical addresses by the TLB when memory is referenced. For memory management purposes, the virtual address range can also be viewed as a set of pages. A segment therefore contains 2^{18} pages of 16 KB each. A virtual page number, combining the segment and the page number within that segment, is a 2^{38} value. Address translations are made at the page level.

The first segments in the virtual address space are special in that they also represent the physical address range. A virtual address in that range bypasses the TLB and accesses physical memory directly.

A TLB can optionally support different page sizes in addition to the architected 16 KB page size. Supported sizes are 16 bytes, 256 KB, 4 MB, and 64 MB. It is very common for the operating system and data to be mapped in consecutive physical pages, which can then be covered by a larger page size.

Security and Protection Support

Twin-64 implements a protection model along two dimensions. The first dimension is **access control** and privilege-level checking. Each page is associated with a page type. Page types include read-only, readwrite, execute, and gateway. There are two privilege levels: user and supervisor mode. Access type and privilege level together form the access control information, which is checked for each instruction and data memory access.

The second dimension of protection is the **protection ID**, which is the segment ID itself. Twin-64 allows a set of segment IDs to be recorded in processor control registers. For each access to a segment that has the protection check bit set, one of the protection ID control registers must match the segment ID. If not, a protection violation trap is raised.

Privilege Changes

Instructions can only access data or execute instructions when the privilege level matches the required privilege level. Two or four levels are the most common implementations. Changing between levels is often done with a kind of trap operation to higher-privilege software, for example the operating system kernel. However, traps are expensive, as they cause the CPU pipeline to be flushed when entering a trap handler.

Twin-64 implements gateways for privilege promotion. A special option on the unconditional branch instruction raises the privilege level when the instruction is executed on a gateway page,

setting the privilege level to that of the gateway page. Returning from a higher privilege level to a code page with lower privilege level automatically resets the lower privilege level. A branch to a higher-privilege page that is not a gateway page results in a privilege violation trap.

Debugging Support

A fundamental requirement is the ability to set instruction and data breakpoints. An instruction breakpoint is triggered by encountering the **TRAP** instruction. The break instruction causes an instruction break trap and passes control to the trap handler. Since break instructions are normally not present in the instruction stream, they must be inserted dynamically in place of the instruction normally found at that address. Upon continuation, if the breakpoint is still valid, it must be reinserted after the original instruction is executed. To facilitate this mechanism, a bit in the status word provides an easy way to trap again after the instruction has been executed.

Data breakpoints are traps taken when a data reference to a page occurs. They do not modify the instruction stream. Depending on the data access, the trap may occur before or after the instruction that accesses the data. A write operation is trapped before the data is written, while a read instruction is trapped after the original instruction has executed.

Input/Output

Twin-64 features a memory-mapped I/O architecture. A portion of the physical address space is dedicated to the I/O subsystem. Within this address range, I/O modules respond to memory load and store instructions directed to their assigned addresses. When implementing a driver, the standard page-level protection mechanism is applied during address translation. A driver, typically implemented to run in privileged mode, can also grant a user program direct control by mapping a user-level I/O page into virtual space without compromising system integrity.

Interrupt System

Interrupts and exceptions are events that occur asynchronously to instruction execution. Depending on the type, they occur either during the execution of an instruction or between instructions. An example of the former is an arithmetic overflow trap; an example of the latter is an external interrupt. Depending on the interrupt or exception type, the instruction is either restarted or execution continues with the next instruction.

Twin-64 implements a simple interrupt system. Each processor features an interrupt input register and an interrupt mask register. When an I/O module receives a command, it also receives information about which interrupt bit to set and which target module ID to send the interrupt to upon completion of the request. The interrupt data is compared to the interrupt mask register. If the particular bit is enabled and interrupts are globally enabled, the target processor enters the interrupt handler.

Since raising an interrupt is simply a matter of storing the interrupt request into the processors interrupt register, I/O modules as well as processors themselves can raise interrupts on a target

processor. Processors can also send interrupts to their peers, enabling low-level processor-to-processor communication.

Instruction Set

Twin-64 features a simple and efficient instruction set. All instructions are of fixed length, specifically a 32-bit word. The total number of instructions is relatively small; however, many instructions encode execution options that make them highly versatile. Great care is taken to ensure that these options do not significantly complicate the pipeline or increase the overall data path length. Instead of supporting a wide variety of addressing modes, as commonly found in CISC-style CPUs, Twin-64 provides two virtual memory addressing modes based on a simple base-offset calculation. No indirection or addressing mode that requires reading a memory operand to compute an address is part of the architecture.

The instruction set is divided into five groups.

- **Memory reference instructions:** Load and store operations for virtual and physical memory.
- **Immediate instructions:** Building immediate values up to 64 bits.
- **Control flow instructions:** Conditional and unconditional branches.
- **Computational instructions:** Arithmetic, Boolean, and bit operations such as extract and deposit.
- **System control instructions:** Managing hardware resources such as caches and TLBs, register movement, traps, and interrupt handling.

Computational Instructions

Arithmetic, logical, and bit-field instructions form the computational class. Twin-64 supports only 64-bit signed arithmetic operations. Any numeric overflow results in a trap.

Binary operations allow negation of either input operand as well as the output. Operand negation provides an efficient mechanism for building complementary bit masks. Bit-field instructions perform extraction and deposit operations, as well as double-register shifts. These serve as the basis for all shift and rotate operations.

Immediate Instructions

Several instructions include bit fields for representing immediate values within the instruction word. Fixed-length instruction word architectures, however, do not allow embedding a full 64-bit immediate in a single instruction. Instead, Twin-64 uses a sequence of two or three instructions to assemble a complete value from parts.

Immediate instructions also support 32-bit address arithmetic. Address calculations wrap around within the 32-bit offset field and never overflow into the segment identifier portion.

Memory Reference Instructions

Memory reference instructions transfer data between memory and general registers. Supported transfer sizes are word, half-word, and byte. In this sense, the architecture resembles a typical load/store design. However, unlike pure load/store architectures, Twin-64 allows many instructions with an operand field to directly access memory. No instruction, however, both reads and writes data memory in the same operation.

With the exception of load-reserved and store-conditional, memory reference instructions support base register modification. When enabled, the base register used for address computation is updated before or after the memory access. A negative offset is applied before address computation, while a positive offset is applied afterward. Indexed addressing further extends the reach by scaling the index according to the data type.

Control Flow Instructions

Control flow instructions are divided into unconditional and conditional branches. Unconditional branches may optionally save the return point in a general register. Returning from a subroutine is accomplished by branching to the saved register value.

Conditional branch instructions compare two operands and transfer control if the condition is satisfied. This eliminates the need for a separate compare instruction followed by a branch.

Branches are further divided into instruction-address-relative and global-address-relative types. The first computes the target relative to the current instruction address, while the second uses a global virtual address as the base.

System Control Instructions

System control instructions manage processor resources such as the TLB, cache, and control registers. Instructions exist for direct access to control registers. The TLB can be managed through insert and remove operations. Cache management includes instructions to flush and purge data. Additional instructions generate traps or provide diagnostic functions for controlling and examining processor state. Most instructions in this group require execution in privileged mode.

Processing Resources

This chapter describes the processing resources and data types in a Twin-64 system.

General Registers

Twin-64 features a set of 16 general purposes registers. They are used for all computations.

	63	0
GR0	Permanent zero	
GR1	Target for ADDIL or General use	
GR2	General use	
	...	
GR14	General use	
GR15	General use	

Figure 1: General registers

Some general registers have dedicated purposes. Register zero always returns a zero value when read, and writing to it has no effect. Although the CPU has only 16 general registers, implementing such a register greatly simplifies instruction design, for example by providing a target when the result is not needed. General register one is a scratch register and also serves as an implicit target for certain instructions. Other general registers may have dedicated roles as defined by the runtime architecture conventions.

Control Registers

Twin-64 has 16 defined control registers, numbered **CR0** to **CR15**. Two instructions allow moving data between a control register and a general register in either direction. A move to a control register always transfers the full 64-bit word. A move from a control register to a general register transfers the value right-aligned to the size of the control register. The following table shows the control registers defined for Twin-64.

	63	32 31										0			
CR0	Processor configuration (PCR)														
CR1	Recovery Counter (RCTR)														
CR2	Reserved														
CR3	Reserved														
CR4	Protection Id 1 (PID1)					W	Protection Id 2 (PID2)					W			
CR5	Protection Id 3 (PID3)					W	Protection Id 4 (PID4)					W			
CR6	Protection Id 5 (PID5)					W	Protection Id 6 (PID6)					W			
CR7	Protection Id 7 (PID7)					W	Protection Id 8 (PID8)					W			
CR8	0					IVA					0				
CR9	External Interrupt Enable Mask (EIEM)														
CR10	Interrupt PSW (IPSW)														
CR11	Interrupt Instruction (IARG0)														
CR12	Interrupt Address (IARG1)														
CR13	Scratch Reg 0														
CR14	Scratch Reg 1														
CR15	Scratch Reg 2														

Figure 2: Control registers

Processor Configuration Register

The processor configuration register contains information about the current system and processor hardware setup. Some of its fields are read-only and set directly by the hardware, while others can be modified by processor-dependent code routines.

63																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																					
----	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Figure 3: Processor Configuration Register

Table 1: Processor Configuration Register Fields

Field	Purpose
Memory Size	The available physical memory in 4Gbyte increments. The minimum size is 4Gbyte, encoded as a zero. The maximum size is 64Gbyte encoded as a 15.
Flags	e.g default endian, etc.
Architecture Id	...
Processor Id Id	...
Processor Core Num	...
Cycles per millisecond	...
...	...

Recovery Counter

decrements when the PSW status bit is set.

Protection Id Registers

Twin-64 allows to record a set of up to 8 protection Id values in the processor control registers CR4 to CR7 which are accessible to the currently executing process. The protection Id is identical to a segment Id. For each access to a segment that has the protection check bit set, one of the protection Id control registers must match the segment Id. If not, a protection violation trap is raised. The rightmost bit of each of the eight PIDs is the write disable (W) bit. When the bit is set, that PID cannot be used to grant write access. See also the chapter on addressing and access control.

Interrupt Vector Address

The Interruption Vector Address (IVA) control register, CR8, holds the base address of an array of service routines assigned to the different interruption classes. The lower 14 bits of the IVA are reserved and must be set to zero, meaning any address written to it must be aligned to a page boundary. The IVA is always located in the first segment, segment 0. The array of interruption service routines is indexed by interruption numbers, as described in the chapter on interrupts.

External Interrupt Enable Mask

The External Interrupt Enable Mask (EIEM) is stored in CR15). It is a 64-bit register, with each bit corresponding to one of the 64 external interrupts. A cleared bit in the EIEM masks the external interrupt associated with that position.

Interruption PSW

Upon an interrupt, the address of the interrupted instruction and the processor status bits are saved to this control register.

Interruption Argument Registers

The Interruption Argument Registers (CR11 and CR12) are used to pass an instruction and a virtual address to an interruption handler. The values of these registers for each interruption class are described in the chapter on interrupts.

Scratch Registers

Control registers CR13, CR14, and CCR15 are scratch registers used by interrupt handlers or similar routines. CR13 and CR14 are privileged and can only be accessed from code running at a privileged level, whereas CR15 is accessible from any privilege level.

Program State Register

The processor state register contains the virtual address of the current instruction along with the processor status flags.




Figure 4: Program State Register

Instruction Address

Processor Status

Table 2: Status Field Bits

Bit	Purpose
M	High priority machine check. When set, high-priority machine checks are masked. This bit is set after an HPMC and cleared after all other interruptions.
E	Little endian access enable. When set, memory access is in little endian format, else big endian format. Endian support is an optional feature. If not implemented, all access is in big endian format.

Continued on next page

Bit	Purpose
R	Recovery counter enable. When set...
X	When set, the processor runs in privileged mode.
T	When set, taken branch traps are enabled. Any branch taken will result in a trap.
H	When set, a higher-privilege transfer trap occurs whenever the following instruction is of a higher privilege.
L	When set, a lower-privilege transfer trap occurs whenever the following instruction is of a lower privilege.
I	When set, external interrupts are enabled.
V	Reserved.
P	Protection identifier validation enable. When set instruction and data references are checked for valid protection identifiers (PIDs). The PRB instruction checks for valid PIDs without causing a trap.
D	Data memory break disable. The D-bit is cleared after the execution of each instruction, except for the RFI instruction which may set it. When set, data memory break traps are disabled. This bit allows a simple mechanism to trap on a data store and then proceed past the trapping instruction.
...	...

Data Types

The fundamental data type is a 64-bit signed integer, with the most significant bit on the left and the least significant bit on the right. All computations are performed on 64-bit signed integers. Memory accesses can be performed at the granularity of a doubleword, word, halfword, or byte. Data loaded as a word, halfword, or byte is sign-extended to 64 bits. Store operations write the corresponding value to memory without modification.




Figure 5: Data types

All memory addresses are expressed as bytes addresses. Address computation uses a 32-bit unsigned math, i.e. numbers wrap around in a 32-bit space and never cross segment boundaries.

Byte Ordering

Twin-64 is a big-endian machine. The following figure shows the memory access for load operations.




Figure 6: Big endian mode

The processor status register reserves a bit position, "E", for future support of mixed big and little endian memory access. Currently this bit is always set to zero.

Translation Lookaside Buffer

Virtual address translations are stored in a translation lookaside buffer (TLB). The TLB is essentially a cache of page table entries representing the virtual pages currently mapped to physical addresses. Depending on the hardware implementation, there may be a single combined TLB or separate instruction and data TLBs. Each TLB entry stores the virtual page number (VPN) along with its attributes and the corresponding physical page number (PPN). The current architecture supports a 36-bit physical address space, allowing a maximum of 64GB of memory.




Figure 7: Tlb info

To reduce TLB processing overhead, the TLB is capable of storing virtual address ranges. There is support for a large page size of 16Kb, 256 Kb, 4Mb and 64 Mb.

Table 3: TLB Fields

Field	Purpose
VPN	The virtual page number.
PPN	The physical page number.
Access Rights	Encodes the access type (read, write, and execute) and the required privilege level for the operation.
Page Size	2-bit field encoding the supported page size: 0 = 16 KByte, 1 = 256 KByte, 2 = 4 MByte, 3 = 64 MByte.
V	Indicates if the entry is valid.
L	When set, the TLB entry is locked and will not be selected during replacement.
U	When set, the page represented by the TLB entry is not cached.
P	When set, the page has protection checking enabled.
M	When set, the page has been modified.
B	When set, a branch instruction executed from this page will cause a branch-taken trap.

The first segments, segment 0 to 15, are special in that they bypass the TLB entirely. A virtual address in the range 0x00_0000_0000 to 0x03_FFFF_FFFF directly maps to the physical address of the same range. The address range represents physical memory and can only be accessed in privileged mode.

Caches

Twin-64 is a cache-visible architecture that supports both unified and separate instruction and data caches. While the hardware directly manages cache line insertion and replacement, software has explicit control over flushing and purging cache line entries. Twin-64 caches are virtually indexed and physically tagged. To accelerate memory access, the cache uses the virtual address to index into the arrays in parallel with the TLB lookup. A cache hit occurs when the physical address tag in the cache matches the physical page address from the TLB.

The architecture does not mandate specific cache line sizes, cache organization, or hierarchy models. Instead, it defines instructions for cache operations, such as deletion or flushing of cache lines.

Each TLB entry indicates whether the corresponding data should be cached. For segments 0 to 15, where the TLB is bypassed, the hardware determines the caching behavior. Memory address ranges are always cached, whereas the I/O address range is not cached.

System Memory

Virtual Address Space

Twin-64 features a 52-bit virtual address range, organized into a 20-bit segment identifier and a 32-bit offset within that segment. The following figure illustrates the structure of a virtual address and how the segment is selected.




Figure 8: Virtual Memory Address Range

A virtual address is translated into a physical address. For translation purposes, the virtual address range can also be viewed as a set of virtual pages. The architected page size is 16 KB, with a 38-bit virtual page number and a 14-bit page offset.

Physical Address Space

Twin-64 architecture features a 36-bit physical memory address range which allows a maximum physical memory of up to 64 GBytes.




Figure 9: Physical Memory Address Range

Segment IDs 015 represent the physical address space, with a maximum size of 64 GB. These segments are directly mapped to physical addresses, without TLB translation or access-rights checking. Only privileged-mode code can access these segments. Control Register 0 contains a field that encodes the actual memory size installed in the system.

Twin-64 uses a memory-mapped I/O architecture. The I/O address space occupies the top 256 MB of the first segment. I/O modules allocate their registers and storage within this range, and access to this area is uncached.

The remaining address space is dedicated to virtual memory. Virtual storage is allocated dynamically, and the TLB provides the necessary translation to physical addresses while enforcing access rights.

Addressing and Access Control

The CPU generates 52-bit virtual addresses. Each virtual memory access must be translated and checked for valid access rights. This chapter describes the address translation process, as well as the access-rights and protection mechanisms.

Address Translation

Each virtual memory access generated by the processor must be translated into a physical address. The translation process maps a virtual page to a physical page. Virtual pages in the first 16 segments bypass the TLB, meaning that the virtual address is identical to the physical address.




Figure 10: Address translation

Access Control

Twin-64 implements a protection model along two dimensions. The first dimension is **access control** and **privilege level**. Each page is associated with a page type: read-only, read-write, execute, or gateway. The page type is encoded in the AccType field, where 0 represents read-only access, 1 represents read-write access, 2 represents execute access, and 3 represents gateway access.

Each memory access is checked against the allowed access type defined in the page descriptor. There are two privilege levels: user mode and supervisor mode. The combination of access type and privilege level forms the access control information, which is verified for every instruction and data memory access.

ID into one of the protection ID registers for the stack owner process, only the corresponding user task is allowed to access its stack data segment with a matching protection ID.

Control Flow

In Twin-64, control flow normally passes to the next sequential instruction in memory unless altered by branch instructions or interruptions. To the programmer, the CPU appears to execute instructions in the order in which they are stored in memory. However, hardware implementations may internally execute instructions out of order. This chapter provides a logical view of the execution model. The sections on branching and interruptions describe how control flow can be altered during program execution.

Unconditional Branches

Unconditional branch instructions alter control flow, and may optionally save the return point to a general register. The branch target address can be computed in two modes:

- **Instruction-relative branches:** The target address is computed by adding an offset to the current instruction address.
- **Segment-relative branches:** The target address is computed relative to the start of the segment, which is address zero.

Both methods use 32-bit unsigned arithmetic, with wrap-around on overflow. For example, suppose the current instruction is at address `0xFFFF_FFFC` and the branch offset is `0x10`. The target address is:

$$0xFFFF_FFFC + 0x10 = 0x1_0000_000C$$

Because address arithmetic is performed modulo 2^{32} , the high-order carry is discarded. The effective target is therefore `0x0000_000C`, and execution continues at that address.

Conditional Branches

Conditional branch instructions combine a comparison or test operation with an Instruction-relative branch. If the condition is met, control transfers to the target instruction. Otherwise, execution continues with the next sequential instruction.

Linkage

Certain branch instructions provide a return path for procedure calls. The return point is the instruction following the branch instruction. The linkage mechanism applies to both intra-segment and inter-segment branches. The return point (segment and offset) is saved in the specified target general register.

Branching and Segments

A branch address consists of a segment identifier and an offset:

- **Intra-segment branches:** Unconditional branches within the same segment. Unconditional branches are always within the same segment.
- **Inter-segment branches:** Unconditional branches into another segment, performed by setting the target address segment and offset portion.

Privilege Level Changes

Branch instructions may change the privilege level. Instructions may only access data or execute code when the current privilege level matches the required level. Switching between levels is often implemented using a trap operation to higher-privileged software such as the operating system kernel. However, traps are expensive, as they typically flush the CPU pipeline when invoking a trap handler.

Twin-64 implements a gateway mechanism for privilege changes. An unconditional branch instruction with the **.G** gateway option raises the privilege level when executed on a gateway page. The branch target continues execution with the privilege level specified by the gateway page, and the privilege bit in the PSW is set. Returning from a higher privilege level to code in a lower-privilege page automatically lowers the current privilege level.

Each instruction fetch compares the privilege level of the code page with the PSW P-bit:

- If the page requires a higher privilege level than the PSW, a privilege violation trap occurs.
- If the page requires a lower privilege level than the PSW, the CPU automatically demotes execution to the lower level.

Traps Associated with Branches

Branch instructions may cause traps based on PSW bits. If the PSW T-bit is set and a branch is taken, a *taken-branch trap* occurs. This mechanism is primarily intended for debugging.

Interruptions

Twin-64 features a simple but powerful interrupt system. Interruptions in Twin-64 are divided into four classes: faults, traps, interrupts, and checks that may occur during instruction execution.

- **Fault** - The current instruction cannot be carried out due to a system problem, such as the absence of a page from main memory. After the system problem has been corrected, the faulting instruction should execute as intended. Faults are synchronous with respect to the instruction stream.
- **Trap** - There are two kind of traps. The types are traps raised because the function requested by the current instruction cannot or should not be carried out. A good example is the arithmetic signed overflow trap. Such an instruction is normally not re-executed. The second type of traps are raised so that the system can intervene before the or after the instruction is executed. An example for this type of trap is the debugging breakpoint trap. Traps are synchronous with respect to the instruction stream.
- **Interrupt** - An external entity such as an I/O module requires attention. Interrupts are asynchronous with respect to the instruction stream.
- **Check** - The processor has detected an internal fault or malfunction. Checks can be either synchronous or asynchronous with respect to the instruction stream.

All four classes of interruptions are handled in the same way.

Interrupt Handling

Handling of an interrupt is implemented as a fast context switch. When an interruption occurs, the hardware saves the PSW at the time of the interrupt to the IPSW control register. PSW in effect at the time of the interruption is saved in the IPSW. Depending on the interrupt type, this can be the current instruction or the next instruction. All status bits, which are part of the PSW, are saved too. The hardware will then set the PSW status bits as follows. The M bit is set if the interruption is a high priority machine check. All other bits are set to zero. Depending on the trap type, additional data is saved to the interrupt argument control registers IARG0 and IARG1.

The control register IVA holds the virtual address of the interrupt vector table. Execution of the interrupt handler begins at the address computed using the trap number:

$$\text{IVA} + (32 * \text{trap number})$$

Each entry consists of 8 instructions that can be executed. It is the responsibility of interruption handlers to unmask external interrupts as soon as possible, so as to minimize interrupt latency. The interruptions are categorized into four groups.

Example. Suppose the interrupt vector table base address is:

$$\text{IVA} = 0\text{x}0000_1000$$

Now, if a **trap number 5** (Illegal Instruction) occurs, the handler entry point is calculated as:

$$0\text{x}0000_1000 + (32 * 5) = 0\text{x}0000_10A0$$

Execution of the interrupt handler starts at **0x0000_10A0**.

Interrupt Vector Table

Table 5: Trap Table

Trap	Name	Instruction Adr	Arg0	Arg1
1	Machine Check	current instruction	check type	-
2	Power Failure	current instruction	-	-
3	Recovery Counter	current instruction	-	-
4	External Interrupt	next instruction	-	-
5	Illegal instruction	current instruction	instruction	-
6	Privileged operation	current instruction	instruction	-
7	Privileged register	current instruction	instruction	-
8	Integer Overflow	current instruction	instruction	-
9	Instruction TLB miss	current instruction	-	-
10	Non-access Instruction TLB miss	current instruction	-	-
11	Instruction memory protection	current instruction	-	-
12	Data TLB miss	current instruction	instruction	target address
13	Non-access Data TLB miss	current instruction	instruction	target address
14	Data memory access rights	current instruction	instruction	target address
15	Data memory protection	current instruction	instruction	target address
16	Page reference	current instruction	instruction	target address
17	Alignment	current instruction	instruction	target address
18	Break Instruction	current instruction	instruction	-
19	Taken branch	next instruction	instruction	-

Instruction Set Overview

This chapter presents the Twin-64 instruction set. Following the typical RISC design, it features a fixed-length format with a small number of regular instruction types. Instructions are grouped into four categories: arithmetic and logical, load and store, branch, and special and system instructions.

Instruction Formats

Twin-64 employs a small number of instruction formats, which are classified according to the number of register operands and the length of the immediate value field. The **signed immediate S19 / special** format provides one register operand and a 19-bit signed immediate field. This format is typically used for branch instructions.

Grp	OpCode	RegR	Opt1	S-IMM-19/special
2	4	4	3	19

The **signed immediate S15 / special** instruction format provides two register operands and a 15-bit signed immediate field. This format is used both by conditional branch instructions and by instructions that operate on a register and an immediate value.

Grp	OpCode	RegR	Opt1	RegB	S-IMM-15/special
2	4	4	3	4	15

The **signed immediate S13 / special** instruction format provides two registers, an additional option field, and a 13-bit signed immediate field. Some instructions use the S-IMM-13 field for special encodings instead of a signed value. This format is used by memory reference instructions, where the address is formed by a base register plus a 13-bit signed offset.

Grp	OpCode	RegR	Opt1	RegB	Opt2	S-IMM-13/special
2	4	4	3	4	2	13

The **register/signed immediate S9 / special** instruction format provides three registers, an additional option field, and a 9-bit signed immediate field. Some instructions use the S-IMM-9 field for special encodings instead of a signed value. This instruction format is used for all computational instructions that require three registers.

Grp	OpCode	RegR	Opt1	RegB	Opt2	RegA	S-IMM-9/special
2	4	4	3	4	2	4	9

The **unsigned immediate U20** format is used by the immediate instruction group to provide a 20-bit unsigned value. This value is then placed at specific positions in the target register Reg R.

Grp	OpCode	RegR	0	U-IMM-20/special
2	4	4	2	20

Arithmetic Instructions

The arithmetic instructions operate on two signed 64-bit operands and store the result in the target register. There are two instruction layouts: signed immediate S15 and register-based. The **immediate operand instruction format** adds the sign-extended 15-bit value to **RegB**. The **register instruction format** uses two registers as input operands, performs the operation, and stores the result in a third register.

The **ADD** and **SUB** instructions perform signed 64-bit integer arithmetic, with numeric overflow triggering an overflow trap. There are no operations on unsigned quantities. The **SHLxA** and **SHRxA** instructions perform an arithmetic left or right shift on the input operand and add another operand to the shifted result, storing the final value in the target register. The **CMP** instruction compares two registers for equality, inequality, less than, and less than or equal conditions, storing the result in the target register.

Bit Operations Instructions

The bit operation instructions fall into two groups. The **AND**, **OR**, and **XOR** instructions perform the logical operations indicated by their names. Like the arithmetic instructions, they have both an immediate instruction format and a register instruction format.

The second group implements bit manipulation operations. The **EXTR** instruction extracts a bit field from a register and places it in another register; optionally, the extracted field can be sign-extended. The **DEP** instruction deposits a bit field into a target register. The **DSR** instruction concatenates two general-purpose registers, shifts the resulting 128-bit quantity to the right, and stores the lower 64 bits in the target register.

Immediate Instructions

The **LDI** and **ADDIL** instructions are used for forming large constants. With a 32-bit instruction word, it is not possible to encode even a full 32-bit constant in a single instruction. The **LDI** instruction includes a qualifier to load a constant value at a defined position in the target register. The **opt** field allows selection of the specific fields in the target register where the constant value is stored. Using a short sequence of instructions, any 64-bit value can be loaded.

The **ADDIL** instruction adds the upper portion of a 32-bit immediate value to the lower 32-bit half of a general register. This instruction is primarily used for address arithmetic, enabling access to any location within a segment. The **LDO** instruction loads an offset into a general register. The address is formed by adding the signed immediate value to a base register. Address arithmetic is performed as unsigned 32-bit arithmetic, wrapping around within the 32-bit range.

Memory Reference Instructions

Twin-64 is a register-memory architecture. It supports the common load/store instructions as well as instructions that combine an operation with a memory reference. Memory reference instructions feature two addressing modes. The first is the base register plus offset mode, and the second is the base register plus register index mode.

The **indexed instruction format** loads the content of the memory address formed by adding the scaled 13-bit immediate field to the base register **RegB**, and adds the fetched data value to **RegA**. The **dw** field specifies the data width: a **dw** value of 0 loads an 8-bit value, 1 loads a 16-bit value, 2 loads a 32-bit value, and 3 loads a 64-bit value. In all cases, the fetched value is sign-extended to 64 bits. Additionally, the **dw** field scales the offset by left-shifting the immediate field according to the data width. For example, a **dw** value of 3 loads a double word from an address formed by shifting the 13-bit offset by three bits before adding it to the base register.

The **register-indexed instruction format** loads the content of the memory address formed by adding **RegX** to the base register **RegB**. The **dw** field specifies the data width, as in the indexed instruction format. The offset in **RegX** is interpreted as a byte offset, independent of the **dw** value. In both instruction formats, the fetched data is sign-extended to 64 bits.

The **LD** and **ST** instructions are the standard load and store instructions. In addition to the two addressing modes described above, they also support base register modification. Depending on the sign of the offset, the base register is updated either before or after computing the effective address. See the instruction descriptions for more details.

The **LDR** and **STC** instructions provide the base support for a pipeline-friendly method for semaphore operations. They only support the indexed instruction mode with double-word access, and the base register modification option is not available.

The **ADD**, **SUB**, **AND**, **OR**, **XOR**, and **CMP** instructions perform the respective operations as described above, with one operand fetched from memory using the indexed addressing modes.

Control Flow Instructions

Control flow is implemented through a set of branch instructions, which can be classified into unconditional and conditional branches.

Unconditional branches fetch the next instruction address from the computed branch target. The **B** and **BR** instructions perform an instruction-relative branch. Both instructions can optionally store the return address in a general register.

The **BV** instruction performs a vectored branch. The content of a general register is added to a base register to form the branch target.

The **BB**, **CBR**, **MBR**, and **ABR** instructions implement conditional branching. If the specified condition is met, a branch to the target address is performed. The target address is always a local address, with the offset being instruction-address relative. Conditional branches adopt a static prediction model where forward branches are assumed not taken, while backward branches are assumed taken.

All branch address arithmetic is performed as 32-bit unsigned arithmetic with wraparound. Adding branch offsets will not overflow into the segment Id portion of the address.

System Instructions

The final instruction group is the **special/system** instruction group. The system control instructions manage CPU resources such as the TLB, cache, and other hardware components. Most instructions in this group require the processor to run in privileged mode.

The **MFCR** and **MTCR** instructions are used to transfer data to and from a control register. Access to some control registers requires privileged mode. The **MFIA** instruction accesses the current instruction address offset, which is important for position-independent code. The **RSM** and **SSM** instructions allow setting or clearing an individual accessible bit in the status portion of the processor state.

The **LDPA**, **PRBR**, and **PRBW** instructions are used to obtain information about the virtual memory system. The **LDPA** instruction returns the physical address of the virtual page, if present in memory. The **PRBx** instructions test whether the desired access mode to an address is allowed.

The translation lookaside buffers and caches are managed by the **ITLB**, **PTLB**, **FCA**, and **PCA** instructions. The **ITLB** and **PTLB** instructions insert or remove translations from the TLB. The **FCA** and **PCA** instructions manage the instruction and data caches, providing options to flush or purge a cache line. There is no instruction to insert data into a cache, as cache insertion is fully controlled by hardware.

The **RFI** instruction restores the processor state from the control registers. Optionally, general registers stored in reserved control registers will also be restored.

The **DIAG** instruction issues hardware-specific implementation commands. An example is the memory barrier command, which ensures that all memory access operations complete before subsequent instructions execute.

The **TRAP** instruction raises a software exception and is typically used to implement a debugging breakpoint.

Instruction Set Descriptions

This chapter describes the Twin-64 instruction set. Each instruction description includes the instruction word layout, a brief description of the instruction, and a high-level pseudo code representation of its operation. The pseudo code routines used in the instruction descriptions are listed in the appendix. Each instruction operation is described using C-style pseudo code. The pseudo code references registers by their class and index. For example, **GR**[5] denotes general register 5, and **CR**[1] denotes control register 1.

Some registers also have dedicated names. For example, the shift amount control register is labelled **SAR**. A subfield of an instruction is selected by adding a dot and the field name in square brackets. For instance, the interrupt enable bit in the PSW register is described as **PSW**.**[I]**.

Additional options for an instruction are specified by a dot followed by a sequence of characters in the assembler syntax. Each character represents a defined option for the instruction. For example, the **AND** instruction can negate its result; in assembler notation, this is written as **AND.N** The order of the individual characters does not matter. All defined options are listed as a character set.

Reserved fields in an instruction word are reserved for future enhancements and should be filled with zero. The instruction explicitly shows the numeric value of these fields. Any other value in a reserved field will result in undefined behavior.

ABR - Add and branch

Syntax

ABR.<cond> RegR, RegB, target

where cond is EQ, LT, NE or LE.

Format

The ABR instruction uses the following format:

2	7	Reg R	Opt 1	Reg B	S-IMM-15/special
2	4	4	3	4	15

Description

The ABR instruction ...

Operation

```
RegR <- RegR + RegB;  
if ( cond( RegR )) IA <- Ia + ofs;  
else IA <- Ia + 4;
```

Exceptions

None.

Notes

None.

ADD Integer Addition

Syntax

ADD RegR, RegB, RegA

ADD RegR, RegB, val

ADD[.D/W/H/B] RegR, ofs (RegB)

ADD[.D/W/H/B] RegR, RegX (RegB)

Format

The ADD instruction uses the register, the immediate-15, the indexed and the register indexed instruction formats.

0	1	RegR	0	RegB	0	RegA	0
2	4	4	3	4	2	4	9

0	1	RegR	1	RegB	val
2	4	4	3	4	15

1	1	RegR	0	RegB	dw	ofs
2	4	4	3	4	2	13

1	1	RegR	1	RegB	dw	RegX	0
2	4	4	3	4	2	4	9

Description

The ADD instruction adds two signed 64-bit operands and stores the result to the target register **RegR**. The instruction supports the immediate, the register and the two memory reference instruction formats. An arithmetic overflow results in a numeric overflow trap. The indexed instruction formats may cause a memory reference associated trap.

Operation

RegR $\leftarrow$ RegB + RegA (register format)

RegR $\leftarrow$ RegB + immOperand() (immediate format)

RegR $\leftarrow$ RegR + memOperand() (indexed formats)

Exceptions

Integer Overflow Trap

Data Alignment Trap

Memory Access Violation Trap

Memory Protection Violation Trap

Data TLB Miss Trap

Notes

None.

ADDIL Add immediate left

Syntax `ADDIL RegR, val`

Format The ADDIL instruction uses the immediate-20 instruction format.

0	9	Reg R	0	val
2	4	4	2	20

Description The ADDIL instruction ...

Operation `RegR <- RegR + ( val << 12 );`

Exceptions None.

Notes None.

AND Boolean Operation

Syntax

```
AND[.C/N] RegR, RegB, RegA
AND[.C/N] RegR, RegB, SignedImmed15

AND[.C/N/D/W/H/B] RegR, SignedImmed13( RegB )
AND[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The AND instruction uses the register, the immediate-19, the indexed and the register indexed instruction formats.

0	3	RegR	N	C	0	RegB	RegA	0	
2	4	4	1	1	1	4	2	4	9

0	3	RegR	N	C	1	RegB	val		
2	4	4	1	1	1	4	15		

1	3	RegR	N	C	0	RegB	dw	ofs	
2	4	4	1	1	1	4	2	13	

1	3	RegR	N	C	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description

The AND instruction performs the binary AND operation on two 64-bit operands and stores the result to the target register **RegR**. Upon input the general register **regB** can be negated. The result can also be negated by setting the "C" bit. The result can be negated by setting the "N" bit. The instruction supports the immediate, the register and the two memory reference instruction formats. The boolean operation itself does not cause any traps. The indexed formats may cause a memory reference associated trap.

Operation

```
if ( instr.[C] ) RegB <- not( RegB );
RegR <- RegB & RegA; (register format)
RegR <- RegB & immOperand(); (immediate format)
RegR <- RegR & memOperand(); (indexed formats)
if ( instr.[N] ) RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

B branch unconditional

Syntax B[.G] target [, RegR]

Format The B instruction uses the immediate-19 format.

2	1	Reg R	0	G	target
2	4	4	2	3	19

Description The B instruction ...

Operation RegR <- addOfs32(PSW.[IA], 4); PSW.[IA] <- addOfs32(
PSW.[IA], target);

Exceptions Instruction Address Trap Instruction Privilege Violation Trap Instruc-
tion TLB Miss Trap

Notes None.

BB branch on bit

Syntax

BB [.T/F] ofs, pos
BB [.T/F] ofs, SAR

Format

The BB instruction uses the the immediate-9 instruction formats

2	4	Reg R	0	A	V	pos	ofs
2	4	4	1	1	1	6	9

Description

The BB instruction ...

Operation

```
if ( instr.[A] ) tmp <- SAR;  
else ( tmp <- instr.[pos];  
if ( checkBit( RegR, instr.[V] ) PSW.[IA] <- addOfs(  
PSW.[IA], ofs );  
else PSW.[IA] <- addOfs( PSW.[IA], 4 );
```

Exceptions

tbd...

Notes

None.

BR branch register

Syntax BR RegB [, RegR]

Format The BR instruction uses the immediate-15 instruction format.

2	2	Reg R	0	Reg B	0
2	4	4	3	4	15

Description The BR instruction ...

Operation `RegR <- addOfs32( PSW.[IA], 4 );      PSW.[IA] <- addOfs32( RegB, target );`

Exceptions None.

Notes None.

BV branch vectored

Syntax

```
RegR <- addOfs32( PSW.[IA], 4 );  
PSW.[IA] <- addOfs32( RegB, RegX );
```

Format

The BV instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description

The BV instruction ...

The indexed formats may cause a memory reference type trap.

Operation

...

Exceptions

...

Notes

None.

CBR compare and branch

Syntax

CBR.<cond> RegR, RegB, target

where cond is EQ, LT, NE or LE.

Format

The CBR instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description

The CBR instruction ...

Operation

...

Exceptions

None.

Notes

None.

CMP Integer Compare

Syntax


```
CMP.<cond> RegR, RegB, RegA
CMP.<cond> RegR, RegB, SignedImmed15

CMP.<cond>[.C/N/D/W/H/B] RegR, SignedImmed13( RegB )
CMP.<cond>[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

where cond_i is EQ, LT, NE or LE.

Format

The CMP instruction uses the register, the immediate-15, the indexed and the register indexed instruction formats.



Description

The **CMP** instruction performs an integer comparison of the two 64bit signed operands and depending on the comparison stores a 0 or 1 in the target register. The indexed formats may cause a memory reference associated trap.

The condition field encodes the comparison options. The options are a 0 for equal, a 1 for less than, a 2 for not equal and a 3 for less than or equal.

Operation

```
cond <- instr.[cond];
RegR <- cmpOp( cond, RegB, RegA ) (register format)
RegR <- cmpOp( cond, RegB, immOperand( )) (immediate format)
RegR <- cmpOp( cond, RegR, memOperand( )) (indexed formats)
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

DEP Deposit Bit Field

Syntax

```
DEP[.Z/I] RegR, RegB, pos, len  
DEP[.Z/I] RegR, RegB, SAR, len
```

Format

The DEP instruction uses the special-13 instruction format with special fields.

0	7	Reg R	0	Reg B	I	A	Z	pos	len
2	4	4	3	4	3			4	4

Description

The DEP instruction deposits the bit field with the specified length in general register RegB into the general register RegR at the specified position. The position field specifies the rightmost bit for the bit field to deposit. The length field specifies the bit size of the field to deposit. The Z bit clears the target register before storing the bit field. If the A bit is set, the shift amount control register is used for obtaining the position value. The I bit specifies that the RegA field contains an immediate value instead of a register.

Operation

```
if ( instr.[Z] ) RegR <- 0;  
if ( instr.[A] ) pos <- shamr;  
else pos <- instr.[pos];  
if ( instr.[I] )  
  RegR <- depositImm( RegR, RegB, pos, instr.[6:0] );  
else  
  RegR <- depositField( RegR, RegB, pos, instr.[6:0] );
```

Exceptions

None

Notes

None.

DIAG - ...

Syntax

DIAG ...

Format

The DIAG instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description

The DIAG instruction ...

The indexed formats may cause a memory reference type trap.

Operation

...

Exceptions

...

Notes

None.

DSR Double Shift Right

Syntax

DSR RegR, RegB, RegA, amt
DSR RegR, RegB, RegA, SAR

Format

The DSR instruction uses the special-9 instruction format.

0	7	Reg R	3	Reg B	0	A	Reg A	0	amt
2	4	4	3	4	2	4	2	6	

Description

The DSR instruction concatenates the general registers RegB and RegA and performs a right shift operation. Register RegB is the left part, register RegA the right part. The lower 64 bits of the result are stored in the general register RegR. The amount field contains the shift amount. If the A bit is set, the shift amount is taken from the shift amount control register.

Operation

```
if ( instr.[A] ) tmp <- shamReg;  
else tmp <- instr.[amt];  
RegR <- doubleShiftR( RegB, RegA, tmp );
```

Exceptions

None.

Notes

None.

EXTR Extract Bit Field

Syntax

```
EXTR[.S] RegR, RegB, pos, len  
EXTR[.S] RegR, RegB, SAR, len
```

Format

The EXTR instruction uses the special-9 instruction format with special fields.

0	7	Reg R	0	Reg B	0	A	S	pos	len
2	4	4	3	4	2			4	9

Description

The EXTR instruction performs a bit field extract specified by the position and length instruction field from general register RegB. The position field specifies the rightmost bit of the bitfield to extract. The length field specifies the bit size if the field to extract. The extracted bit field is stored right justified in the general register RegR. If the S bit is set, the result is sign extended. If the A bit is set, the shift amount control register is used for obtaining the position value instead of the length instruction field.

Operation

```
if ( instr.[A] tmp <- shamt;  
else tmp <- instr.[pos];  
RegR <- extractField( RegB, pos, len );  
if ( instr.[S] ) RegR <- signExtend( RegR, pos );
```

Exceptions

None

Notes

None.

FCA - flush from cache

Syntax FCA RegR, RegB

Format The PTLB instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The FCA instruction ...

Operation ...

Exceptions None.

Notes None.

ITLB insert TLB entry

Syntax ITLB ...;

Format The ITLB instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description The ITLB instruction ...

The indexed formats may cause a memory reference type trap.

Operation ...

Exceptions ...

Notes None.

LD Load register

Syntax

LD[.D/W/H/B] RegR, SignedImmed13(RegB)
LD.D/W/H/B] RegR, RegX (RegB)

Format

The LD instruction uses the indexed and the register indexed instruction formats.

1	7	Reg R	0	M	0	Reg B	dw	S-IMM-13
2	4	4	1	1	1	4	2	13

1	7	Reg R	0	M	1	Reg B	dw	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description

The LD instruction ...

The indexed formats may cause a memory reference type trap.

Operation

RegR <- memLoad(memOperand()) (indexed formats)

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

LDI Load immediate

Syntax LDI [.L/S/U] RegR, val

Format The LDI instruction uses the immediate-20 instruction format.

0	9	Reg R	opt	val
2	4	4	2	20

Description The LDI instruction ...

Operation

$$\text{RegR} \leftarrow ((\text{val} \& 0\text{x}\text{FFFFF}) \ll 12); \text{ (.L option)}$$
$$\text{RegR} \leftarrow ((\text{val} \& 0\text{x}\text{FFFFF}) \ll 32); \text{ (.S option)}$$
$$\text{RegR} \leftarrow ((\text{val} \& 0\text{x}\text{FFF}) \ll 52); \text{ (.U option)}$$

Exceptions None.

Notes None.

LDO load offset

Syntax LDO RegR, ofs (RegB)

Format The LDO instruction uses the immediate-15 instruction format.

0	10	Reg R	0	Reg B	ofs
2	4	4	3	4	15

Description The LDO instruction ...

Operation $\text{RegR} \leftarrow \text{RegR} + \text{ofs};$

Exceptions None.

Notes None.

LDR Load register reserved

Syntax LDR RegR, SignedImmed13(RegB)

Format The LDR instruction uses the indexed instruction format.

1	9	Reg R	0	Reg B	3	S-IMM-13
2	4	4	0	4	2	13

Description The LDR instruction ...

The indexed formats may cause a memory reference type trap.

Operation ??? fix ...

RegR <- memLoad(memOperand()) (indexed formats)

Exceptions Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes None.

LPA - load physical address

Syntax LPA ...

Format The BV instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description The LPA instruction ...

The indexed formats may cause a memory reference type trap.

Operation ...

Exceptions ...

Notes None.

MBR move and branch

Syntax

MBR.<cond> RegR, RegB, target

where jcond_i is EQ, LT, NE or LE.

Format

The MBR instruction uses the xxx instruction format.

0	6	Reg R	cond	0	Reg B	ofs
2	4	4	2	1	4	15

Description

The MBR instruction ...

Operation

...

Exceptions

None.

Notes

None.

MFIA - Move from instruction address register

Syntax MFIA RegR

Format The MFIA instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The MTCR instruction ...

Operation ...

Exceptions None.

Notes None.

MFCR Move from Control register

Syntax MFCR RegR, CRegB

Format The MFCR instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The CBR instruction ...

Operation ...

Exceptions None.

Notes None.

MTCR Move to Control register

Syntax MTCR RegR, CRegB

Format The MFCR instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The MTCR instruction ...

Operation ...

Exceptions None.

Notes None.

NOP - No operation

Syntax NOP]

Format The RFI instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The NOP instruction ...

Operation ...

Exceptions None.

Notes None.

OR Boolean Operation

Syntax

```
OR[.C/N] RegR, RegB, RegA
OR[.C/N] RegR, RegB, SignedImmed15

OR[.C/N/D/W/H/B] RegR, SignedImmed13( RegB )
OR[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The OR instruction uses the register, the immediate-19, the indexed and the register indexed instruction formats.

0	4	RegR	N	C	0	RegB	RegA	0	0
2	4	4	1	1	1	4	2	4	9

0	4	RegR	N	C	1	RegB	val
2	4	4	1	1	1	4	15

1	4	RegR	N	C	0	RegB	dw	ofs
2	4	4	1	1	1	4	2	13

1	4	RegR	N	C	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description

The **OR** instruction performs the binary OR operation on two 64-bit operands and stores the result to the target register **RegR**. Upon input the general register **regB** can be negated. The result can also be negated by setting the "C" bit. The result can be negated by setting the "N" bit. The instruction supports the immediate, the register and the two memory reference instruction formats. The boolean operation itself does not cause any traps. The indexed formats may cause a memory reference associated trap.

Operation

```
if ( instr.[C] RegB <- not( RegB );
RegR <- RegB | RegA (register format)
RegR <- RegB | immOperand( ) (immediate format)
RegR <- RegR | memOperand( ) (indexed formats)
if ( instr.[N] RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

PCA - purge from cache

Syntax PCA RegR, RegB

Format The PTLB instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The PCA instruction ...

Operation ...

Exceptions None.

Notes None.

PRB - probe virtual address

Syntax PRB ...

Format The PRB instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description The PRB instruction ...

The indexed formats may cause a memory reference type trap.

Operation ...

Exceptions ...

Notes None.

PTLB - purge from TLB

Syntax PTLB RegR, RegB

Format The PTLB instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The PTLB instruction ...

Operation ...

Exceptions None.

Notes None.

RFI - Return from interrupt

Syntax RFI [RegR]

Format The RFI instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The RFI instruction ...

Operation ...

Exceptions None.

Notes None.

RSM - Reset system mask

Syntax RSM RegR,

Format The MFIA instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The RSM instruction ...

Operation ...

Exceptions None.

Notes None.

SHLxA Shift left and add

Syntax

SHLxA RegR, RegB, RegA
SHLxA RegR, RegB, SignedImmed13

Format

The SHLxA instruction uses the register and the immediate instruction formats.

0	8	Reg R	amt	0	Reg B	0	Reg A	0
2	4	4	1 2		4	2	4	9

1	8	Reg R	amt	0	Reg B	2	S-IMM-13
2	4	4	1 2		4	2	13

Description

The SHLxA instruction ...

Operation

RegR $\leftarrow$ (RegB $\ll$ amt) + RegA (register format)
RegR $\leftarrow$ (RegB $\ll$ amt) + immOperand() (immediate format)

Exceptions

Integer Overflow Trap

Notes

None.

SHRxA Shift right and add

Syntax

SHRxA RegR, RegB, RegA
SHRxA RegR, RegB, SignedImmed13

Format

The SHRxA instruction uses the register and immediate instruction formats.

0	8	Reg R	amt	0	Reg B	0	Reg A	0
2	4	4	1 2		4	2	4	9

1	8	Reg R	amt	0	Reg B	2	S-IMM-13
2	4	4	1 2		4	2	13

Description

The SHRxA instruction ...

Operation

RegR <-(RegB >> amt) + RegA (register format)
RegR <- (RegB >> amt) + immOperand() (immediate format)

Exceptions

Integer Overflow Trap

Notes

None.

SSM - Set system mask

Syntax SSM RegR,

Format The MFIA instruction uses the xxx instruction format.

2	5	Reg R	cond	0	Reg B	ofs
2	4	4	3		4	15

Description The SSM instruction ...

Operation ...

Exceptions None.

Notes None.

ST Store register

Syntax

ST[.D/W/H/B] RegR, SignedImmed13(RegB)
ST.D/W/H/B] RegR, RegX (RegB)

Format

The ST instruction uses the indexed and the register indexed instruction formats.

1	8	Reg R	0	M	0	Reg B	dw	S-IMM-13
2	4	4	1	1	1	4	2	13

1	8	Reg R	0	M	1	Reg B	dw	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description

The ST instruction ...

The indexed formats may cause a memory reference type trap.

Operation

RegR <- memLoad(memOperand()) (indexed formats)

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

STC Store register conditional

Syntax STC RegR, SignedImmed13(RegB)

Format The STC instruction uses the indexed instruction format.

1	8	Reg R	0	Reg B	3	S-IMM-13
2	4	4	0	4	2	13

Description The ST instruction ...

The indexed formats may cause a memory reference type trap.

Operation RegR <- memLoad(memOperand()) (indexed formats)

Exceptions

- Data Alignment Trap
- Memory Access Violation Trap
- Memory Protection Violation Trap
- Data TLB Miss Trap

Notes None.

SUB Integer Subtraction

Syntax

```
SUB RegR, RegB, RegA
SUB RegR, RegB, SignedImmed15

SUB[.D/W/H/B] RegR, SignedImmed13( RegB )
SUB[.D/W/H/B] RegR, RegX ( RegB )
```

Format

The SUB instruction uses the register, the immediate, the indexed and the register indexed instruction formats.

0	2	RegR	0	RegB	0	RegA	0
2	4	4	3	4	2	4	9

0	2	RegR	1	RegB	val		
2	4	4	3	4	15		

1	2	RegR	0	RegB	dw	ofs	
2	4	4	3	4	2	13	

1	2	RegR	1	RegB	dw	RegX	0
2	4	4	3	4	2	4	9

Description

The SUB instruction subtracts a signed 64-bit operand from a general register **RegB** and stores the result to the target register **RegR**. The instruction supports the immediate, the register and the two memory reference instruction formats. An arithmetic overflow results in a numeric overflow trap. The indexed formats may cause a memory reference associated trap.

Operation

```
RegR <- RegB - RegA (register format)
RegR <- RegB - immOperand( ) (immediate format)
RegR <- RegR - memOperand( ) (indexed formats)
```

Exceptions

Integer Overflow Trap
Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

TRAP - ...

Syntax

TRAP ...

Format

The TRAP instruction uses the XXX instruction formats.

2	7	Reg R	0	M	1	Reg B	0	Reg X	0
2	4	4	1	1	1	4	2	4	9

Description

The TRAP instruction ...

The indexed formats may cause a memory reference type trap.

Operation

...

Exceptions

...

Notes

None.

XOR Boolean Operation

Syntax

```
XOR[.N] RegR, RegB, RegA
XOR[.N] RegR, RegB, SignedImmed15

OR[.N/D/W/H/B] RegR, SignedImmed13( RegB )
OR[.N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The XOR instruction uses the register, the immediate, the indexed and the register indexed instruction formats.

0	4	Reg R	N	0	0	Reg B	0	Reg A	0
2	4	4	3			4	2	4	9

0	4	Reg R	N	0	1	Reg B	S-IMM-15
2	4	4	3			4	15

1	4	Reg R	N	0	0	Reg B	dw	S-IMM-13
2	4	4	3			4	2	13

1	4	Reg R	N	0	1	Reg B	dw	Reg X	0
2	4	4	3			4	2	4	9

Description

The XOR instruction performs the binary XOR operation on two 64-bit operands and stores the result to the target register **RegR**. The result can be negated by setting the "N" bit. The instruction supports the immediate, the register and the two memory reference instruction formats. The boolean operation itself does not cause any traps. The indexed formats may cause a memory reference associated trap.

Operation

```
RegR <- RegB ^ RegA (register format)
RegR <- RegB ^ immOperand( ) (immediate format)
RegR <- RegR ^ memOperand( ) (indexed formats)
if ( instr.[N] RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

Input/Output

Concepts

Twin-64 features memory a mapped I/O architecture.

memory mapped

dedicated address space in physical memory

bus concept

module concept

modules have registers - load / store, however registers do not have to be 64 bit wide.

the system bus features up to 64 modules. that is: 64 HPA pages, need: 1Mb address range.

a module listens to three address ranges

- hard physical address range. a page. privileged.

- soft physical address range. allocated by the module for further space. aligned on a page boundary.

- broadcast address range. Mirrors the system bus range. A write to a register in that range applies to all corresponding registers of a module on the same bus.

modules are processors, memory and I/O cards.

we could keep processor stats in the HPA space

memory module has entire memory as SPA space

I/O Address Space

As shown in the overall physical memory layout, the physical memory address range dedicated to I/O modules is located from 0x00_F000_0000 to 0x00_FFFF_FFFF. The following figure depicts the I/O memory address space ranges in more details




Figure 13: I/O Memory Address Range

The I/O address range is a 256Mbyte area in the first 4 Gbyte segment, i.e. segment 0. The area is divided into several ranges. The first range contains the processor dependent code. There is a set of library functions to manage the processor itself but also provide low-level primitives for the boot process and the operating system.

Hard physical address range

The **hard physical address** range contains a I/O page for each I/O modules installed. There room for 64 I/O modules.

Broadcast address range

The **broadcast address** range mirrors the hard physical address range. However, a write to these locations is a write operation to which all I/O modules will react.

Soft physical address range

The **soft physical address** range is an area where space is allocated to an I/O module. The I/O module reacts to a memory access operation in its configured soft physical address range.

I/O Elements

I/O elements are 128 bytes, i.e. 16 regs

I/O elements can also be two pages, one priv and one user level, mapped...

Appendix - Instruction Commentary

xxx

Appendix - Instruction Notation Routines

describe the routines used in the instruction set operation part...

Table 6: Instruction Notation Routines

Function	Description
...	...
...	...

Appendix -TLB and Caches

xxx

