

TWIN-64

Architecture Document

Helmut Fieres
October 28, 2025

Contents

Tables of Contents	iv
List of Figures	v
List of Tables	vii
 I Introduction	 1
Introduction	3
Twin-64	3
Design Goals	3
This book	4
 II Architecture	 5
System Organization	7
A RegisterMemory Architecture	7
System	7
Processor	8
Virtual Memory	8
Security and Protection Support	8
Privilege Changes	9
Debugging Support	9
Input/Output	9
Interrupt System	10
 Processing Resources	 11
General Registers	11
Control Registers	12
Processor State Register	14
Data Types	15
Byte Ordering and Alignment	16
Translation Lookaside Buffer	16
Caches	18
Instruction Set	18
 System Memory	 21
Virtual Address Space	21
Address Space Layout	22
 Addressing and Access Control	 23

CONTENTS

Address Translation	23
Access Control	23
Access Protection	24
Access Check Flow	25
Control Flow	27
Unconditional Branches	27
Conditional Branches	27
Linkage	27
Inter-segment Branches	28
Privilege Level Changes	28
Traps Associated with Branches	28
Interruptions	29
Interrupt Handling	29
Interrupt Vector Table	30
 III Instruction Set	 31
Instruction Set	33
Arithmetic Instructions	33
Bit Operations Instructions	34
Immediate Instructions	34
Memory Reference Instructions	34
Control Flow Instructions	35
System Instructions	35
Instruction Formats	36
 Instruction Set Descriptions	 39
ABR - Add and branch	40
ADD - Integer Addition	41
ADDIL - Add immediate left	42
AND - Boolean Operation	43
B Branch Unconditional	44
BB Branch on Bit	45
BE - Branch External	46
BR Branch Register	47
BV Branch Vectored	48
CBR - Compare and Branch	49
CMP - Compare	50
DEP - Deposit Bit Field	52
DIAG - Diagnose Instruction	53
DSR Double Shift Right	54
EXTR - Extract Bit Field	55
FxCA - flush from cache	56
ITLB Insert TLB Entry	57
LD Load Register	58
LDIL Load Immediate Left	59
LDO Load Offset	60
LDR Load Register Reserved	61
LPA Load Physical Address	62

CONTENTS

MBR - Move and Branch	63
MFIA - Move from Instruction Address	64
MFCR Move From Control Register	65
MTCR Move To Control Register	66
NOP No Operation	67
OR Boolean Operation	68
PCA - Purge from Cache	69
PRB Probe Virtual Address	70
PxTLB Purge TLB Entry	71
RFI Return from Interrupt	72
RSM Reset System Mask	73
SHLxA Shift Left and Add	74
SHRxA - Shift Right and Add	75
SSM Set System Mask	76
ST Store Register	77
STC Store Register Conditional	78
SUB - Integer Subtraction	79
TRAP Trap Instruction	80
XOR Boolean Operation	81
Synthetic Instructions	83
Immediate Operations	83
Register Operations	83
Arithmetic Operations	84
Shift and Rotate Operations	84
Memory Reference Operations	84
 IV I/O Subsystem	 85
Input/Output	87
Concepts	87
I/O Address Space	87
Hard physical address range	88
Broadcast address range	89
Soft physical address range	89
I/O Elements	89
 V Runtime Environment	 91
Runtime Overview	93
System Environment	93
Runtime Environment	93
Register Usage	94
Task Execution	94
 Stack Layout	 95
Stack Frame Marker	95
Stack Frame Marker	95
 Calling Convention	 97
Local Call	97

CONTENTS

Linkage Table	97
Inter module calls	97
 VI Processor Design	 99
 VII Simulator Environment	 101
 VIII Appendix	 103
Appendix - Instruction Commentary	105
Comparison and Branch Condition Encoding	105
Appendix - Instruction Notation Routines	109
Appendix - Programming Notes	111
xxx	111
Appendix - TLB and Caches	113
Appendix - Diagnostic Functions	115
xxx	115
Appendix - Glossary	117

List of Figures

1	Software Accessible Registers	11
2	General registers	11
3	Control registers	12
4	Processor Configuration Register	13
5	Program State Register	14
6	Data types	16
7	Big endian mode	16
8	Tlb info	17
9	Virtual Memory Address Range	21
10	Memory Address Ranges	22
11	Address translation	23
12	Access control information field	24
13	Protection Id	24
14	Address translation	25
15	I/O Memory Address Range	88
16	Runtime Overview	93
17	Stack Frame	95
18	Stack Frame Marker	96
19	Processor Configuration Register	97

LIST OF FIGURES

List of Tables

1	Processor Configuration Register Fields	13
2	Status Field Bits	15
3	TLB Fields	17
4	Access type checking	24
5	Trap Table	30
6	Instruction Notation Routines	109
7	Diagnostic Functions	115

Part I

Introduction

Introduction

Designers of CPUs in the seventies and early eighties almost all used a microcoded approach with hundreds of instructions. RISC was just beginning to emerge. Registers were typically not fully general-purpose but often had special functions or meanings, and many instructions were rather complicated, though designers believed them useful. Compiler writers, however, often used only a small subset, ignoring these complex instructions because they were too specialized. In the late eighties and nineties the shift moved to RISC-based designs, with large sets of registers, fixed instruction word lengths, large virtual memory address ranges, and instructions that were in general much more pipeline-friendly. Today, processors are highly complex, with superscalar deep pipelines and multilevel cache hierarchies—just a few of the hallmarks of modern CPU design. CPU designs have become so large and complex that no single person can completely understand them. But what if there were a simple design that one person could fully understand?

Twin-64

Welcome to Twin-64. Twin-64 is a simple 64-bit CPU with a register-memory model and segmented virtual memory. The design is heavily influenced by Hewlett-Packard's PA-RISC architecture, which was initially a 32-bit RISC-style load/store machine. Almost all of the key architectural features come from there. However, other processors such as the Data General MV8000, the DEC Alpha, and the IBM PowerPC were also influential or at least investigated for their design choices. In addition, the Blitz-64 architecture was studied. Blitz-64, a design developed at Portland State University, offered an interesting approach in that it only supported 64-bit signed arithmetic, avoiding common errors from mixing signed and unsigned arithmetic.

Where does the name Twin-64 come from? Obviously, the 64 indicates that the CPU is a 64-bit CPU. The Twin part will become clearer when implementations of the CPU are described in later chapters. One implementation represents a dual-issue instruction design with two ALUs for computation. Nevertheless, a very basic implementation could just offer single-issue, single-ALU execution as well.

Design Goals

While it is not a design goal to build an industry-strength, competitive CPU, the CPU should nevertheless be useful and largely follow established design practices. For example, instructions should not be invented just because they seem useful, but rather designed with the assembler and compilers in mind. Instruction set design is CPU design. Register-memory architectures were typically microcoded, complex-instruction machines. In contrast, Twin-64 instructions will be hardwired and are in general pipeline-friendly, avoiding data and control hazards and stalls where possible.

This book

This document serves as a comprehensive guide to the Twin-64 processor, its instruction set, and runtime environment. It is organized into several parts, each focusing on a key aspect of the system.

Part Two introduces the Twin-64 architecture, providing an overview of the overall system and processors structure, core components, and design philosophy. This sets the stage for understanding how instructions are executed and how data moves through the system.

Part Three presents the instruction set in detail. It covers computational, immediate, memory reference, branch, and system control instructions. These chapters serve as the definitive reference for programmers and designers alike, offering precise information on encoding, operation, and usage.

Part Four explores the I/O subsystem, detailing the mechanisms by which the processor communicates with external devices and handles input/output operations efficiently.

Part Five discusses the runtime environment and software conventions. It explains program interaction with Twin-64, including calling conventions, exception handling, and runtime support for higher-level languages.

Part Six examines processor design. While multiple design approaches are possible, this section provides a reference implementation that highlights key architectural concepts, practical trade-offs, and performance considerations.

Finally, the **Appendix** provides supplementary material, including additional design notes, tables, and clarifications to support the main text and offer deeper insight into the Twin-64 system.

Part II

Architecture

System Organization

This chapter introduces Twin-64 by highlighting the major components and design concepts.

A RegisterMemory Architecture

Modern RISC CPUs are typically load / store designs and feature a large number of general-purpose registers. Operations take place between registers. Twin-64 follows a slightly different route. It has fewer general registers, and in addition to registerregister computation, a register-memory model is also supported. However, there are no instructions that read from and write to memory in a single instruction cycle, since this would complicate the design considerably.

Twin-64 implements a **register-memory architecture** offering only a few addressing modes for the operand, combined with a fixed instruction word length. For most instructions one operand is a register, while the other is a flexible operand that can be an immediate, another register, or a memory address from which data is obtained or stored. The result is placed in the register, which is also the first operand. These types of machines are also called two-address machines.

In contrast to similar historical register-memory designs, Twin-64 has no operations that both read from and write to memory in the same instruction. For example, there is no increment memory instruction, since this would require two memory operations and is generally not pipeline-friendly. Memory data access is performed at the word, half-word, or byte level. Besides the implicit operand fetch in a computational instruction, there are dedicated load/store instructions. In addition to the registerimmediate and registermemory operand modes, one operand mode supports the three-operand model ($R_s = R_a \text{ OP } R_b$), specifying two source registers and a result register. This allows Twin-64 to support both registermemory and load/store operation styles.

System

A Twin-64 system is organized around a central system bus that interconnects the key components of the architecture. Processors, main memory modules, and I/O modules are all attached as peers to this bus, enabling uniform communication and data exchange. Larger configurations may also include I/O adapters, which bridge slower peripheral buses to the central system bus while preserving a consistent programming model.

Processors initiate memory and I/O transactions via the bus and can also receive external events such as interrupts or system resets through the same interface. This modular structure provides flexibility in system design, allowing configurations to scale from simple processormemory systems to complex multiprocessor environments with diverse I/O subsystems.

Processor

The Twin-64 processor implements the instruction set architecture and provides the computational resources of the system. It contains a general register set, control and status registers, and the program state word. To support virtual memory, each processor includes a Translation Lookaside Buffer (TLB) that translates virtual addresses into physical addresses. Instruction and data caches, while optional in principle, are essential for performance and tightly integrated with the memory interface. Regardless of configuration, all processor-generated addresses are virtual and must be resolved by the TLB before memory access.

Virtual Memory

Twin-64 offers a large global address range, organized in segments. A segment number and offset together form the global virtual address. Segments are also associated with access rights and protection identifiers. The CPU itself can run in user or privileged mode. Twin-64 always generates virtual addresses at the CPU level. The global virtual memory range is a 52-bit space organized into a 20-bit segment ID and a 32-bit byte offset. There are 2^{20} segments with 2^{32} bytes each.

Virtual addresses are translated to physical addresses by the TLB when memory is referenced. For memory management purposes, the virtual address range can also be viewed as a set of pages. A segment therefore contains 2^{18} pages of 16 KB each. A virtual page number, combining the segment and the page number within that segment, is a 2^{38} value. Address translations are made at the page level.

The first segments in the virtual address space are special in that they also represent the physical address range. A virtual address in that range bypasses the TLB and accesses physical memory directly.

A TLB can optionally support different page sizes in addition to the architected 16 KB page size. Supported sizes are 16 bytes, 256 KB, 4 MB, and 64 MB. It is very common for the operating system and data to be mapped in consecutive physical pages, which can then be covered by a larger page size.

Security and Protection Support

Twin-64 implements a protection model along two dimensions. The first dimension is **access type control** and privilege-level checking. Each page is associated with a page type. Page types include read-only, readwrite, execute, and gateway. There are two privilege levels: user and supervisor mode. Access type and privilege level together form the access control information, which is checked for each instruction and data memory access.

The second dimension of protection is the **protection ID**, which is the segment ID itself. Twin-64 allows a set of segment IDs to be recorded in processor control registers. For each access to a segment that has the protection check bit set, one of the protection ID control registers must match the segment ID. If not, a protection violation trap is raised.

Privilege Changes

Instructions can only access data or execute instructions when the privilege level matches the required privilege level. Two or four levels are the most common implementations. Changing between levels is often done with a kind of trap operation to higher-privilege software, for example the operating system kernel. However, traps are expensive, as they cause the CPU pipeline to be flushed when entering a trap handler.

Twin-64 implements gateways for privilege promotion. A special option on the unconditional branch instruction raises the privilege level when the instruction is executed on a gateway page, setting the privilege level to that of the gateway page. Returning from a higher privilege level to a code page with lower privilege level automatically resets the lower privilege level. A branch to a higher-privilege page that is not a gateway page results in a privilege violation trap.

Debugging Support

A fundamental requirement is the ability to set instruction and data breakpoints. An instruction breakpoint is triggered by encountering the **TRAP** instruction. The break instruction causes an instruction break trap and passes control to the trap handler. Since break instructions are normally not present in the instruction stream, they must be inserted dynamically in place of the instruction normally found at that address. Upon continuation, if the breakpoint is still valid, it must be reinserted after the original instruction is executed. To facilitate this mechanism, a bit in the status word provides an easy way to trap again after the instruction has been executed.

Data breakpoints are traps taken when a data reference to a page occurs. They do not modify the instruction stream. Depending on the data access, the trap may occur before or after the instruction that accesses the data. A write operation is trapped before the data is written, while a read instruction is trapped after the original instruction has executed.

Input/Output

Twin-64 implements a memory-mapped I/O architecture. A portion of the physical address space is reserved for I/O modules, which respond to memory load and store instructions directed to their assigned addresses. Standard page-level protection applies during address translation, ensuring that only privileged software (typically device drivers) can safely access I/O regions. User programs may be granted direct access to I/O by mapping a user-level I/O page into virtual memory, without compromising overall system integrity.

Direct I/O is performed using the standard load and store instructions to access I/O modules. These operations are entirely controlled by software and may be executed in user mode if the virtual mapping allows it.

Direct Memory Access (DMA) modules can transfer data between memory and I/O devices independently of the processor. DMA transfers operate on physical memory, requiring that the relevant pages be locked and protected from conflicting access by the operating system. DMA

transactions are initiated by issuing DMA commands, and command chaining is supported to enable complex or continuous data transfers without processor intervention.

Interrupt System

Interrupts and exceptions are events that occur asynchronously to instruction execution. Depending on the type, they occur either during the execution of an instruction or between instructions. An example of the former is an arithmetic overflow trap; an example of the latter is an external interrupt. Depending on the interrupt or exception type, the instruction is either restarted or execution continues with the next instruction.

Twin-64 implements a simple interrupt system. Each processor features an interrupt input register and an interrupt mask register. When an I/O module receives a command, it also receives information about which interrupt bit to set and which target module ID to send the interrupt to upon completion of the request. The interrupt data is compared to the interrupt mask register. If the particular bit is enabled and interrupts are globally enabled, the target processor enters the interrupt handler.

Since raising an interrupt is simply a matter of storing the interrupt request into the processors interrupt register, I/O modules as well as processors themselves can raise interrupts on a target processor. Processors can also send interrupts to their peers, enabling low-level processor-to-processor communication.

Processing Resources

This chapter describes the processing resources and data types in a Twin-64 system. Twin-64 provides a set of general-purpose registers, a set of control registers, and the Program State Word (PSW), which contains the current instruction address and processor status. The following figure presents a high-level overview of these register sets.




Figure 1: Software Accessible Registers

General registers are used for computation and address storage, control registers manage processor state, and the PSW holds the current instruction address along with the processor status.

General Registers

Twin-64 features a set of 16 general purposes registers. They are used for all computations including address arithmetic.

	63	0
GR0	Permanent zero	
GR1	Target for ADDIL or General use	
GR2	General use	
	...	
GR14	General use	
GR15	General use	

Figure 2: General registers

Some general registers have dedicated purposes. Register zero always returns a zero value when read, and writing to it has no effect. Although the CPU has only 16 general registers,

implementing such a register greatly simplifies instruction design, for example by providing a target when the result is not needed. General register one is a scratch register and also serves as an implicit target for certain instructions. Other general registers may have dedicated roles as defined by the runtime architecture conventions.

Control Registers

Twin-64 has 16 defined control registers, numbered **CR0** to **CR15**. Two instructions allow moving data between a control register and a general register in either direction. A move to a control register always transfers the full 64-bit word. A move from a control register to a general register transfers the value right-aligned to the size of the control register. The following table shows the control registers defined for Twin-64.

	63	32	31	0
CR0	Processor configuration (PCR)			
CR1	Recovery Counter (RCTR)			
CR2	Reserved			
CR3	Reserved			
CR4	Protection Id 1 (PID1)	W	Protection Id 2 (PID2)	W
CR5	Protection Id 3 (PID3)	W	Protection Id 4 (PID4)	W
CR6	Protection Id 5 (PID5)	W	Protection Id 6 (PID6)	W
CR7	Protection Id 7 (PID7)	W	Protection Id 8 (PID8)	W
CR8	Interrupt Vector Table Address (IVA)			
CR9	External Interrupt Enable Mask (EIEM)			
CR10	Interrupt PSW (IPSW)			
CR11	Interrupt Instruction (IARG0)			
CR12	Interrupt Address (IARG1)			
CR13	Scratch Reg 0			
CR14	Scratch Reg 1			
CR15	Scratch Reg 2			

Figure 3: Control registers

Processor Configuration Register

The processor configuration register contains information about the current system and processor hardware setup. Some of its fields are read-only and set directly by the hardware, while others can be modified by processor-dependent code routines.




Figure 4: Processor Configuration Register

Table 1: Processor Configuration Register Fields

Field	Purpose
Memory Size	The available physical memory in 4Gbyte increments. The minimum size is 4Gbyte, encoded as a zero. The maximum size is 64Gbyte encoded as a 15.
Flags	e.g default endian, etc.
Architecture Id	...
Processor Id Id	...
Processor Core Num	...
...	...
Cycles per millisecond	...

Recovery Counter

When enabled, the recovery counter decrements on each instruction cycle. When the leftmost bit becomes a one, a recovery trap is scheduled. The R bit in the processor status word is set, the recovery trap is enabled.

Protection Id Registers

Twin-64 allows to record a set of up to 8 protection Id values in the processor control registers CR4 to CR7 which are accessible to the currently executing process. The protection Id is identical to a segment Id. For each access to a segment that has the protection check bit set, one of the protection Id control registers must match the segment Id. If not, a protection violation trap is raised. The rightmost bit of each of the eight PIDs is the write disable (W) bit. When the bit is set, that PID cannot be used to grant write access. See also the chapter on addressing and access control.

Interrupt Vector Table Address

The Interrupt Vector Table Address (IVA) control register, CR8, holds the base address of an array of service routines assigned to the different interruption classes. The lower 14 bits of the IVA are reserved and must be set to zero, meaning any address written to it must be aligned to a page boundary. The IVA is always located in the first segment, segment 0. The array of interruption service routines is indexed by interruption numbers, as described in the chapter on interrupts.

External Interrupt Enable Mask

The External Interrupt Enable Mask (EIEM) is stored in CR15). It is a 64-bit register, with each bit corresponding to one of the 64 external interrupts. A cleared bit in the EIEM masks the external interrupt associated with that position.

Interruption PSW

Upon an interrupt, the address of the interrupted instruction and the processor status bits are saved to this control register.

Interruption Argument Registers

The Interruption Argument Registers (CR11 and CR12) are used to pass an instruction and a virtual address to an interruption handler. The values of these registers for each interruption class are described in the chapter on interrupts.

Scratch Registers

Control registers CR13, CR14, and CCR15 are scratch registers used by interrupt handlers or similar routines. CR13 and CR14 are privileged and can only be accessed from code running at a privileged level, whereas CR15 is accessible from any privilege level.

Program State Register

The processor state register contains the virtual address of the current instruction along with the processor status flags.




Figure 5: Program State Register

Instruction Address

The instruction address portion of the PSW registers contains the segment id and the byte offset into the segment for the currently executing instruction.

Processor Status

Table 2: Status Field Bits

Bit	Purpose
M	High priority machine check. When set, high-priority machine checks are masked. This bit is set after an HPMC and cleared after all other interruptions.
E	Little endian access enable. When set, memory access is in little endian format, else big endian format. Endian support is an optional feature. If not implemented, all access is in big endian format.
R	Recovery counter enable. When set...
X	When set, the processor runs in privileged mode.
T	When set, taken branch traps are enabled. Any branch taken will result in a trap.
I	When set, external interrupts are enabled.
V	Reserved.
P	Protection identifier validation enable. When set instruction and data references are checked for valid protection identifiers (PIDs). The PRB instruction checks for valid PIDs without causing a trap.
D	Data memory break disable. The D-bit is cleared after the execution of each instruction, except for the RFI instruction which may set it. When set, data memory break traps are disabled. This bit allows a simple mechanism to trap on a data store and then proceed past the trapping instruction.
...	...

Data Types

The fundamental data type is a 64-bit signed integer, with the most significant bit on the left and the least significant bit on the right. All computations are performed on 64-bit signed integers. Memory accesses can be performed at the granularity of a double-word, word, halfword, or byte. Data loaded as a word, halfword, or byte is sign-extended to 64 bits. Store operations write the corresponding value to memory without modification.




Figure 6: Data types

All memory addresses are expressed as bytes addresses. Address computation uses a 32-bit unsigned math, i.e. numbers wrap around in a 32-bit space and never cross segment boundaries.

Byte Ordering and Alignment

Twin-64 is a big-endian machine. The following figure shows the memory access for load operations.




Figure 7: Big endian mode

The processor status register reserves a bit position, "E", for future support of mixed big and little endian memory access. Currently this bit is always set to zero.

Twin-64 enforces natural alignment for memory accesses. Byte, half-word, word, and double-word data types must be aligned to 1, 2, 4, and 8-byte boundaries respectively. If an access does not meet the required alignment, a data-alignment trap is raised.

Translation Lookaside Buffer

Virtual address translations are stored in a translation lookaside buffer (TLB). The TLB is essentially a cache of page table entries representing the virtual pages currently mapped to physical addresses. Depending on the hardware implementation, there may be a single combined TLB or separate instruction and data TLBs. Each TLB entry stores the virtual page number

(VPN) along with its attributes and the corresponding physical page number (PPN). The current architecture supports a 36-bit physical address space, allowing a maximum of 64GB of memory.




Figure 8: Tlb info

To reduce TLB processing overhead, the TLB is capable of storing virtual address ranges. There is support for a large page size of 16Kb, 256 Kb, 4Mb and 64 Mb.

Table 3: TLB Fields

Field	Purpose
VPN	The virtual page number.
PPN	The physical page number.
Access Rights	Encodes the access type (read, write, and execute) and the required privilege level for the operation.
Page Size	2-bit field encoding the supported page size: 0 = 16 KByte, 1 = 256 KByte, 2 = 4 MByte, 3 = 64 MByte.
V	Indicates if the entry is valid.
L	When set, the TLB entry is locked and will not be selected during replacement.
U	When set, the page represented by the TLB entry is not cached.
P	When set, the page has protection checking enabled.
M	When set, the page has been modified. This only applies to data TLBs or unified TLBs.
B	When set, a branch instruction executed from this page will cause a branch-taken trap. This inly applies for instruction TLBs or unified TLBs.

The first segments, segment 0 to 15, are special in that they bypass the TLB entirely. A virtual address in the range 0x00_0000_0000 to 0x03_FFFF_FFFF directly maps to the physical address of the same range. The address range represents physical memory and can only be accessed in privileged mode.

Caches

Twin-64 is a cache-visible architecture that supports both unified and separate instruction and data caches. While the hardware directly manages cache line insertion and replacement, software has explicit control over flushing and purging cache line entries. Twin-64 caches are virtually indexed and physically tagged. To accelerate memory access, the cache uses the virtual address to index into the arrays in parallel with the TLB lookup. A cache hit occurs when the physical address tag in the cache matches the physical page address from the TLB.

The architecture does not mandate specific cache line sizes, cache organization, or hierarchy models. Instead, it defines instructions for cache operations, such as deletion or flushing of cache lines.

Each TLB entry indicates whether the corresponding data should be cached. For segments 0 to 15, where the TLB is bypassed, the hardware determines the caching behavior. Memory address ranges are always cached, whereas the I/O address range is not cached.

Instruction Set

Twin-64 features a simple and efficient instruction set. All instructions are of fixed length, specifically a 32-bit word. The total number of instructions is relatively small, but many include execution options that make them highly versatile. Great care has been taken to ensure these options do not significantly complicate the pipeline or lengthen the data path.

Unlike CISC-style processors that support a wide variety of addressing modes, Twin-64 provides just two virtual memory addressing modes, both based on a simple baseoffset calculation. No indirection or mode requiring a memory operand to compute an address is part of the architecture.

The instruction set is organized into five groups:

- **Computational instructions:** Arithmetic, Boolean, and bit-manipulation operations such as extract and deposit. Twin-64 supports only 64-bit signed arithmetic. Any numeric overflow results in a trap.
- **Immediate instructions:** Instructions for constructing immediate values up to 64 bits. Several instructions embed immediate fields within the instruction word, but a full 64-bit constant requires a sequence of two or three instructions. Immediate instructions also support 32-bit address arithmetic, where calculations wrap around within the 32-bit offset field and never overflow into the segment identifier portion.
- **Memory reference instructions:** Load and store operations for both virtual and physical memory. These transfer data between memory and general registers in sizes of double-word, word, halfword, or byte. The design resembles a load/store architecture, but unlike strict load/store machines, many Twin-64 instructions allow memory operands directly. No instruction, however, both reads and writes data memory in the same operation.
- **Control flow instructions:** Both unconditional and conditional branches are supported. Unconditional branches may optionally save the return address in a general-purpose reg-

ister. Conditional branches compare two operands and transfer control if the condition is satisfied, eliminating the need for a separate compare instruction.

- **System control instructions:** Instructions for managing processor resources such as the TLB, cache, and control registers. Control registers can be accessed directly; the TLB can be updated with insert and remove operations; and the cache can be flushed or purged. Additional instructions generate traps or provide diagnostic functions for examining processor state. Most system instructions require execution in privileged mode.

The book section on the instruction set will describe each instruction group in detail, with formats, descriptions, and examples.

System Memory

Virtual Address Space

Twin-64 features a 52-bit virtual address range, organized into a 20-bit segment identifier and a 32-bit offset within that segment. The following figure illustrates the structure of a virtual address and how the segment is selected.




Figure 9: Virtual Memory Address Range

A virtual address is translated into a physical address. For translation purposes, the virtual address range can also be viewed as a set of virtual pages. The architected page size is 16 KB, with a 38-bit virtual page number and a 14-bit page offset.

Address Space Layout

Twin-64 architecture features a 36-bit physical memory address range which allows a maximum physical memory of up to 64 GBytes.




Figure 10: Memory Address Ranges

Segment IDs 0 to 15 represent the physical address space, with a maximum size of 64 GB. These segments are directly mapped to physical addresses, without TLB translation or access-rights checking. Only privileged-mode code can access these segments. Control Register 0 contains a field that encodes the actual memory size installed in the system.

Twin-64 uses a memory-mapped I/O architecture. The I/O address space occupies the top 256 MB of the first segment. I/O modules allocate their registers and storage within this range, and access to this area is uncached.

The remaining address space is dedicated to virtual memory. Virtual storage is allocated dynamically, and the TLB provides the necessary translation to physical addresses while enforcing access rights.

Addressing and Access Control

The CPU generates 52-bit virtual addresses. Each virtual memory access must be translated and checked for valid access rights. This chapter describes the address translation process, as well as the access-rights and protection mechanisms.

Address Translation

Each virtual memory access generated by the processor must be translated into a physical address. The translation process maps a virtual page to a physical page. Virtual pages in the first 16 segments bypass the TLB, meaning that the virtual address is identical to the physical address.




Figure 11: Address translation

Access Control

Twin-64 implements a protection model along two dimensions. The first dimension is **access control** and **privilege level**. Each page is associated with a page type: read-only, read-write, execute, or gateway. The page type is encoded in the AccType field, where 0 represents read-only access, 1 represents read-write access, 2 represents execute access, and 3 represents gateway access.

Each memory access is checked against the allowed access type defined in the page descriptor. There are two privilege levels: user mode and supervisor mode. The combination of access type and privilege level forms the access control information, which is verified for every instruction and data memory access.




Figure 12: Access control information field

For read access, the privilege level must be at least as high as the PL1 field; for write access, it must be at least as high as PL2. Exceptions are execute and gateway pages, where write access is allowed only for privileged code. For execute access, the privilege level must be at least PL1 and no higher than PL2. If the instructions privilege level falls outside these bounds, a privilege violation trap is raised. If the page type does not match the instructions access type, an access violation trap is raised. In both cases, the instruction is aborted, and a memory protection trap occurs.

Table 4: Access type checking

Type	read	write	execute
0 - read-only	PL <= PL1	not allowed	not allowed
1 - read-write	PL <= PL1	PL <= PL2	not allowed
2 - execute	PL <= PL1	PL == 0	PL2 <= PL <= PL1
3 - gateway	PL <= PL1	PL == 0	PL2 <= PL <= PL1

Access Protection

The second dimension of protection is the **protection ID**. Twin-64 allows the processor to store up to eight protection ID values in the control registers **CR4** through **CR7**, which are accessible to the currently executing process. The protection ID is identical to a segment ID. For each access to a segment with the protection-check bit set, at least one of the protection ID registers must match the segment ID; otherwise, a protection violation trap is raised.



Figure 13: Protection Id

When protection ID checking is enabled, the eight protection identifiers are compared with the segment ID of the target address to validate an access. The rightmost bit of each protection ID register serves as a write-disable (W) bit. When the W bit is set, that protection ID cannot be used to grant write access. If the PSW P-bit is cleared, protection ID and write-disable checks are bypassed.

Protection ID checking is typically used to enforce controlled access to shared or private memory regions. A common example is the stack data segment, which is accessible at user privilege level. Since the CPU implements a global virtual memory address space, any task could potentially access the stack of another task. By enabling protection ID checking and loading the segment

ID into one of the protection ID registers for the stack owner process, only the corresponding user task is allowed to access its stack data segment with a matching protection ID.

Note: Is the protection Id bit in the processor status word required? When we have the rule that privilege code can access anything, and user code will always be subject to protection Id checking, the bit is not needed.

Access Check Flow

The following figures summarizes the protection and access mode checking during address translation.




Figure 14: Address translation

For a successfully found TLB entry, the segment portion of the virtual address is compared against the set of protection Id registers. If there is no match, a protection Id trap is raised. The access mode of the instruction is compared against the access rights and TLB flags of the TLB entry. If the access mode does not match an access rights trap or a privilege violation trap is raised.

Control Flow

In Twin-64, control normally proceeds to the next sequential instruction in memory unless explicitly changed by a branch or interrupted by an exception. From the programmers perspective, instructions execute in program order. Internally, however, the hardware may employ techniques such as out-of-order execution. This chapter presents the logical execution model, describing how branching and interruptions alter control flow.

Unconditional Branches

Unconditional branch instructions always redirect control flow and may optionally save the return address in a general-purpose register. The branch target address can be computed in two ways:

- **Instruction-relative branches:** The target is obtained by adding an offset to the address of the current instruction.
- **Segment-relative branches:** The target is computed relative to the start of the segment (address zero of the segment).

Both methods use 32-bit unsigned arithmetic with wrap-around on overflow. For example, if the current instruction resides at `0xFFFF_FFFC` and the branch offset is `0x10`, the calculation is:

$$0xFFFF_FFFC + 0x10 = 0x1_0000_000C$$

Since the addition is performed modulo 2^{32} , the high-order carry is discarded. The effective branch target is therefore `0x0000_000C`.

Conditional Branches

Conditional branches combine a comparison with an instruction-relative branch. If the specified condition is true, control transfers to the branch target; otherwise, execution continues sequentially. Conditional branches are always instruction-relative.

Linkage

Certain branch instructions support procedure calls by saving the return address (the instruction immediately following the branch) into a general-purpose register. The linkage mechanism applies to both intra-segment and inter-segment branches. The saved return address consists of both the segment identifier and the instruction offset.

Inter-segment Branches

Unconditional inter-segment branches are branches to another segment. The offset is computed from the offset portion in the address specified plus the immediate value encoded in the **BE** instruction.

Privilege Level Changes

Branch instructions may alter the current privilege level. Execution is only permitted if the current privilege level satisfies the requirements of the code or data being accessed. Traditionally, privilege changes are handled through traps into the operating system kernel. While secure, this approach incurs significant overhead due to pipeline flushes and handler setup. Twin-64 provides a more efficient gateway mechanism. An unconditional branch with the **.G** gateway option may enter a gateway page, raising the privilege level. The processor then executes at the privilege level specified by the gateway page, with the PSW privilege bit updated accordingly. When returning to code in a lower-privilege page, the CPU automatically lowers the current privilege level. Each instruction fetch checks the privilege level of the code page against the PSW P-bit:

- If the page requires higher privilege than the PSW indicates, a privilege violation trap occurs.
- If the page requires lower privilege, execution privilege level is automatically demoted.
- If the **G** option on the branch instruction **Bis** set, the execution privilege is set to the privilege level of the gateway page where the branch instruction resides.

Traps Associated with Branches

Branch instructions may trigger traps depending on the PSW state. For example, if the PSW **T**-bit is set and a branch is taken, a *taken-branch trap* is raised. This feature is primarily intended for debugging.

Interruptions

Twin-64 features a simple but powerful interrupt system. Interruptions in Twin-64 are divided into four classes: faults, traps, interrupts, and checks that may occur during instruction execution.

- **Fault** - The current instruction cannot be carried out due to a system problem, such as the absence of a page from main memory. After the system problem has been corrected, the faulting instruction should execute as intended. Faults are synchronous with respect to the instruction stream.
- **Trap** - There are two kind of traps. The types are traps raised because the function requested by the current instruction cannot or should not be carried out. A good example is the arithmetic signed overflow trap. Such an instruction is normally not re-executed. The second type of traps are raised so that the system can intervene before the or after the instruction is executed. An example for this type of trap is the debugging breakpoint trap. Traps are synchronous with respect to the instruction stream.
- **Interrupt** - An external entity such as an I/O module requires attention. Interrupts are asynchronous with respect to the instruction stream.
- **Check** - The processor has detected an internal fault or malfunction. Checks can be either synchronous or asynchronous with respect to the instruction stream.

All four classes of interruptions are handled in the same way.

Interrupt Handling

Handling of an interrupt is implemented as a fast context switch. When an interruption occurs, the hardware saves the PSW at the time of the interrupt to the IPSW control register. PSW in effect at the time of the interruption is saved in the IPSW. Depending on the interrupt type, this can be the current instruction or the next instruction. All status bits, which are part of the PSW, are saved too. The hardware will then set the PSW status bits as follows. The M bit is set if the interruption is a high priority machine check. All other bits are set to zero. Depending on the trap type, additional data is saved to the interrupt argument control registers IARG0 and IARG1.

The control register IVA holds the virtual address of the interrupt vector table. Execution of the interrupt handler begins at the address computed using the trap number:

$$\text{IVA} + (32 * \text{trap number})$$

Each entry consists of 8 instructions that can be executed. It is the responsibility of interruption handlers to unmask external interrupts as soon as possible, so as to minimize interrupt latency. The interruptions are categorized into four groups.

Example. Suppose the interrupt vector table base address is:

$$\text{IVA} = 0\text{x}0000_1000$$

Now, if a **trap number 5** (Illegal Instruction) occurs, the handler entry point is calculated as:

$$0\text{x}0000_1000 + (32 * 5) = 0\text{x}0000_10A0$$

Execution of the interrupt handler starts at 0x0000_10A0.

Interrupt Vector Table

Table 5: Trap Table

Trap	Name	Instruction Adr	Arg0	Arg1
1	Machine Check	current instruction	check type	-
2	Power Failure	current instruction	-	-
3	Recovery Counter	current instruction	-	-
4	External Interrupt	next instruction	-	-
5	Illegal instruction	current instruction	instruction	-
6	Privileged operation	current instruction	instruction	-
7	Privileged register	current instruction	instruction	-
8	Integer Overflow	current instruction	instruction	-
9	Instruction TLB miss	current instruction	-	-
10	Non-access Instruction TLB miss	current instruction	-	-
11	Instruction memory protection	current instruction	-	-
12	Data TLB miss	current instruction	instruction	target address
13	Non-access Data TLB miss	current instruction	instruction	target address
14	Data memory access rights	current instruction	instruction	target address
15	Data memory protection	current instruction	instruction	target address
16	Page reference	current instruction	instruction	target address
17	Alignment	current instruction	instruction	target address
18	Break Instruction	current instruction	instruction	-
19	Taken branch	next instruction	instruction	-

Part III

Instruction Set

Instruction Set

This part of the book presents the Twin-64 instruction set. Following the typical RISC design, it features a fixed-length format with a small number of regular instruction types. The total number of instructions is relatively small; however, many instructions encode execution options that make them highly versatile. Great care is taken to ensure that these options do not significantly complicate the pipeline or increase the overall data path length. Instead of supporting a wide variety of addressing modes, as commonly found in CISC-style CPUs, Twin-64 provides two virtual memory addressing modes based on a simple base-offset calculation. No indirection or addressing mode that requires reading a memory operand to compute an address is part of the architecture.

The instruction set is divided into five groups.

- **Memory reference instructions:** Load and store operations for virtual and physical memory.
- **Immediate instructions:** Building immediate values up to 64 bits.
- **Control flow instructions:** Conditional and unconditional branches.
- **Computational instructions:** Arithmetic, Boolean, and bit operations such as extract and deposit.
- **System control instructions:** Managing hardware resources such as caches and TLBs, register movement, traps, and interrupt handling.

Arithmetic Instructions

The arithmetic instructions operate on two signed 64-bit operands and store the result in the target register. There are two instruction layouts: signed immediate S15 and register-based. The immediate operand instruction format adds the sign-extended 15-bit value to **RegB**. The register instruction format uses two registers as input operands, performs the operation, and stores the result in a third register.

The **ADD** and **SUB** instructions perform signed 64-bit integer arithmetic, with numeric overflow triggering an overflow trap. There are no operations on unsigned quantities. The **SHLxA** and **SHRxA** instructions perform an arithmetic left or right shift on the input operand and add another operand to the shifted result, storing the final value in the target register. The **CMP** instruction compares two registers for equality, inequality, less than, and less than or equal conditions, storing the result in the target register.

Bit Operations Instructions

The bit operation instructions fall into two groups. The **AND**, **OR**, and **XOR** instructions perform the logical operations indicated by their names. Like the arithmetic instructions, they have both an immediate instruction format and a register instruction format.

The second group implements bit manipulation operations. The **EXTR** instruction extracts a bit field from a register and places it in another register; optionally, the extracted field can be sign-extended. The **DEP** instruction deposits a bit field into a target register. The **DSR** instruction concatenates two general-purpose registers, shifts the resulting 128-bit quantity to the right, and stores the lower 64 bits in the target register.

Immediate Instructions

The **LDI** and **ADDIL** instructions are used for forming large constants. With a 32-bit instruction word, it is not possible to encode even a full 32-bit constant in a single instruction. The **LDI** instruction includes a qualifier to load a constant value at a defined position in the target register. The **opt** field allows selection of the specific fields in the target register where the constant value is stored. Using a short sequence of instructions, any 64-bit value can be loaded.

The **ADDIL** instruction adds the upper portion of a 32-bit immediate value to the lower 32-bit half of a general register. This instruction is primarily used for address arithmetic, enabling access to any location within a segment. The **LDO** instruction loads an offset into a general register. The address is formed by adding the signed immediate value to a base register. Address arithmetic is performed as unsigned 32-bit arithmetic, wrapping around within the 32-bit range.

Memory Reference Instructions

Twin-64 is a registermemory architecture. It supports the common load/store instructions as well as instructions that combine an operation with a memory reference. Memory reference instructions feature two addressing modes. The first is the base register plus offset mode, and the second is the base register plus register index mode.

The indexed instruction format loads the content of the memory address formed by adding the scaled 13-bit immediate field to the base register **RegB**, and adds the fetched data value to **RegA**. The **dw** field specifies the data width: a **dw** value of 0 loads an 8-bit value, 1 loads a 16-bit value, 2 loads a 32-bit value, and 3 loads a 64-bit value. In all cases, the fetched value is sign-extended to 64 bits. Additionally, the **dw** field scales the offset by left-shifting the immediate field according to the data width. For example, a **dw** value of 3 loads a double word from an address formed by shifting the 13-bit offset by three bits before adding it to the base register.

The register-indexed instruction format loads the content of the memory address formed by adding **RegX** to the base register **RegB**. The **dw** field specifies the data width, as in the indexed instruction format. The offset in **RegX** is interpreted as a byte offset, independent of the **dw** value. In both instruction formats, the fetched data is sign-extended to 64 bits.

The LD and ST instructions are the standard load and store instructions. In addition to the two addressing modes described above, they also support base register modification. Depending on the sign of the offset, the base register is updated either before or after computing the effective address. See the instruction descriptions for more details.

The LDR and STC instructions provide the base support for a pipeline-friendly method for semaphore operations. They only support the indexed instruction mode with double-word access, and the base register modification option is not available.

The ADD, SUB, AND, OR, XOR, and CMP instructions perform the respective operations as described above, with one operand fetched from memory using the indexed addressing modes.

Control Flow Instructions

Control flow is implemented through a set of branch instructions, which can be classified into unconditional and conditional branches.

Unconditional branches fetch the next instruction address from the computed branch target. The B and BR instructions perform an instruction-relative branch. Both instructions can optionally store the return address in a general register.

The BV instruction performs a vectored branch. The content of a general register is added to a base register to form the branch target.

The BB, CBR, MBR, and ABR instructions implement conditional branching. If the specified condition is met, a branch to the target address is performed. The target address is always a local address, with the offset being instruction-address relative. Conditional branches adopt a static prediction model where forward branches are assumed not taken, while backward branches are assumed taken.

All branch address arithmetic is performed as 32-bit unsigned arithmetic with wraparound. Adding branch offsets will not overflow into the segment Id portion of the address.

System Instructions

The final instruction group is the **special/system** instruction group. The system control instructions manage CPU resources such as the TLB, cache, and other hardware components. Most instructions in this group require the processor to run in privileged mode.

The MFCR and MTCR instructions are used to transfer data to and from a control register. Access to some control registers requires privileged mode. The MFIA instruction accesses the current instruction address offset, which is important for position-independent code. The RSM and SSM instructions allow setting or clearing an individual accessible bit in the status portion of the processor state.

The LDPA, PRBR, and PRBW instructions are used to obtain information about the virtual memory system. The LDPA instruction returns the physical address of the virtual page, if present in memory. The PRBx instructions test whether the desired access mode to an address is allowed.

The translation lookaside buffers and caches are managed by the ITLB, PTLB, FCA, and PCA instructions. The ITLB and PTLB instructions insert or remove translations from the TLB. The FCA and PCA instructions manage the instruction and data caches, providing options to flush or purge a cache line. There is no instruction to insert data into a cache, as cache insertion is fully controlled by hardware.

The RFI instruction restores the processor state from the control registers. Optionally, general registers stored in reserved control registers will also be restored.

The DIAG instruction issues hardware-specific implementation commands. An example is the memory barrier command, which ensures that all memory access operations complete before subsequent instructions execute.

The TRAP instruction raises a software exception and is typically used to implement a debugging breakpoint.

Instruction Formats

Twin-64 employs a small number of instruction formats, which are classified according to the number of register operands and the length of the immediate value field. The **signed immediate S19 / special (S19)** format provides one register operand and a 19-bit signed immediate field. This format is typically used for branch instructions.

Grp	OpCode	RegR	Opt1	S-IMM-19/special
2	4	4	3	19

The **signed immediate S15 / special (S15)** instruction format provides two register operands and a 15-bit signed immediate field. This format is used both by conditional branch instructions and by instructions that operate on a register and an immediate value.

Grp	OpCode	RegR	Opt1	RegB	S-IMM-15/special
2	4	4	3	4	15

The **signed immediate S13 / special (S13)** instruction format provides two registers, an additional option field, and a 13-bit signed immediate field. Some instructions use the S-IMM-13 field for special encodings instead of a signed value. This format is used by memory reference instructions, where the address is formed by a base register plus a 13-bit signed offset.

Grp	OpCode	RegR	Opt1	RegB	Opt2	S-IMM-13/special
2	4	4	3	4	2	13

The **register / signed immediate S9 / special (S9)** instruction format provides three registers, an additional option field, and a 9-bit signed immediate field. Some instructions use the S-IMM-9 field for special encodings instead of a signed value. This instruction format is used for all computational instructions that require three registers.

Grp	OpCode	RegR	Opt1	RegB	Opt2	RegA	S-IMM-9/special
2	4	4	3	4	2	4	9

The **unsigned immediate U20 (U20)** format is used by the immediate instruction group to provide a 20-bit unsigned value. This value is then placed at specific positions in the target register **Reg R**.

Grp	OpCode	RegR	0	U-IMM-20/special
2	4	4	2	20

In the instruction descriptions, each instruction format is referred to by its abbreviation (such as S9, S15, and so on). These formats specify the arrangement of operands, control bits, and optional immediate values within the instruction word. For some instructions, the immediate field directly encodes a constant, while in others the field is reserved or implicitly set according to the instructions encoding rules.

Instruction Set Descriptions

This chapter provides a detailed description of the Twin-64 instruction set. Each instruction entry includes:

- The instruction word layout.
- A concise description of the instruction.
- A high-level pseudo code representation of its operation.

Instruction operations are expressed using C-style pseudo code. The routines and conventions used are listed in the appendix. Registers are referenced by class and index:

- `GR[5]` general register 5
- `CR[1]` control register 1
- Some registers have dedicated names, e.g., the shift amount control register is **SAR**.

Subfields of an instruction are accessed using a dot and the field name in square brackets. For example:

- The interrupt enable bit in the PSW register: `PSW.[I]`.

Instruction options are specified in assembler syntax by a dot followed by one or more characters, each representing a defined option. For example:

- `AND.N . . .` the `AND` instruction with the negate result option.
- The order of option characters does not matter; all defined options are listed in the relevant character set.

Fields marked as reserved are intended for future enhancements and should always be set to zero. Any other value in a reserved field results in undefined behavior.

ABR - Add and branch

Syntax ABR.<cond> RegR, RegB, target

Format The ABR instruction uses the signed immediate S15 format:

2	11	Reg R	cond	Reg B	ofs
2	4	4	3	4	15

Description The ABR instruction adds **RegB** to **RegR** and stores the result in **RegR**. If the condition specified in the **cond** field is satisfied, execution continues at the current instruction plus the offset; otherwise, execution proceeds with the next instruction.

Conditions The **cond** field encodes the branch condition as follows:

Value	Mnemonic	Branch if
0	EQ	RegR == 0
1	LT	RegR < 0
2	NE	RegR != 0
3	GE	RegR >= 0
4	LE	RegR <= 0
5	EV	RegR.[0] == 0
6	GT	RegR > 0
7	OD	RegR.[0] != 0

Operation

```
RegR <- RegR + RegB;
if ( cond( RegR )) psw.[IA] <- psw.[IA] + ofs;
else                psw.[IA] <- psw.[IA] + 4;
```

Exceptions Overflow trap.

Notes The **ABR** instruction extends the **CMP** functionality by integrating a conditional branch. With a mode field of three bits available for the **cond** field, all six standard relational conditions (EQ, NE, LT, LE, GT, GE) are directly encoded, avoiding the need for operand swapping or result inversion.

See also the instruction notes on comparison and branch Condition encoding in the appendix.

ADD Integer Addition

Syntax

```
ADD RegR, RegB, RegA
ADD RegR, RegB, val

ADD[.D/W/H/B] RegR, ofs ( RegB )
ADD[.D/W/H/B] RegR, RegX ( RegB )
```

Formats

The ADD instruction uses the register, the signed immediate S15, the indexed and the register indexed instruction formats.

0	1	RegR	0	RegB	0	RegA	0
2	4	4	3	4	2	4	9

0	1	RegR	1	RegB	val		
2	4	4	3	4	15		

1	1	RegR	0	RegB	dw	ofs	
2	4	4	3	4	2	13	

1	1	RegR	1	RegB	dw	RegX	0
2	4	4	3	4	2	4	9

Description

The ADD instruction adds two signed 64-bit operands and stores the result in the target register **RegR**. This instruction supports the immediate, register, and both memory-reference formats. An arithmetic overflow triggers a numeric overflow trap, and indexed instruction formats may cause a memory-reference related trap.

Operation

RegR <- RegB + RegA;	(S9)
RegR <- RegB + signExt(val);	(S15)
RegR <- RegR + memLoad(RegB + signExt(ofs << dw));	(S13)
RegR <- RegR + memLoad(RegB + RegX);	(S9)

Exceptions

Integer Overflow Trap
Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

ADDIL - Add immediate left

Syntax `ADDIL RegR, val`

Format The ADDIL instruction uses the unsigned immediate U20 instruction format.

0	9	Reg R	0	val
2	4	4	2	20

Description The **ADDIL** instruction adds an immediate value to a general register using 32-bit unsigned arithmetic. The 20-bit immediate value is shifted left by 12 bits (padded with zeros on the right) before being added to **RegR**. The resulting value is stored in the general register **GR1**. Any overflow that occurs during the addition is ignored.

Operation

RegR <- RegR + (val << 12);

Exceptions None.

Notes None.

AND - Boolean Operation

Syntax

```
AND[.C/N] RegR, RegB, RegA
AND[.C/N] RegR, RegB, val

AND[.C/N/D/W/H/B] RegR, ofs( RegB )
AND[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The AND instruction uses the register, the signed immediate S15, the indexed and the register indexed instruction formats.

0	3	RegR	N	C	0	RegB	0	RegA	0
2	4	4	1	1	1	4	2	4	9

0	3	RegR	N	C	1	RegB	val		
2	4	4	1	1	1	4	15		

1	3	RegR	N	C	0	RegB	dw	ofs	
2	4	4	1	1	1	4	2	13	

1	3	RegR	N	C	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description

The AND instruction performs a bitwise AND operation on two 64-bit operands and stores the result in the target register **RegR**. The second operand, **RegB** or the loaded content, can be negated by setting the "C" bit, and the result can be negated by setting the "N" bit. This instruction supports the immediate, register, and both memory-reference formats. The Boolean operation itself does not trigger any traps, but the indexed formats may cause a memory-reference related trap.

Operation

```
tmp <- RegA;                (register format)
tmp <- immOperand( instr );  (immediate format)
tmp <- memOperand( instr );  (indexed formats)

if ( instr.[C] ) tmp <- not( tmp );
RegR <- RegB & tmp;
if ( instr.[N] ) RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

B Branch Unconditional

Syntax `B[.G] target [, RegR ]`

Format The B instruction uses the signed immediate S19 format.

2	1	RegR	0	G	ofs
2	4	4	2	1	19

Description The B instruction performs an unconditional instruction-addressrelative branch. The target address is computed by sign-extending the immediate offset, shifting it left by two bits, and adding the result to the current instruction address. The return link is stored in `RegR`.

The optional `.G` form enables the *gateway branch* option. In this case, the processors privilege level may be raised to the privilege level stored in the TLB entry of the page from which the gateway instruction is fetched. The change is only performed if it results in a higher privilege level; otherwise, the current privilege level remains unchanged. Execution then continues at the computed target address with the (possibly new) privilege level.

Operation

```
??? revise

if ( RegR != 0 ) RegR <- addOfs32( PSW.[IA], 4 );

if ( instr.[G] ) {

    if ( TLB[PSW.[IA]].[X] > PSW.[X] )
        PSW.[X] <- TLB[PSW.[IA]].[X];
}

PSW.[IA] <- addOfs32( PSW.[IA], ofs << 2 );
```

Exceptions Instruction Privilege Violation Trap
Instruction TLB Miss Trap
Taken Branch Trap

Notes None.

BB Branch on Bit

Syntax BB[.T/F] ofs, pos
BB[.T/F] ofs, SAR

Format The BB instruction uses the signed immediate S13 format.

2	8	RegR	0	A	V	pos	ofs
2	4	4	1	1	1	6	13

Description The BB instruction tests a bit in register **RegR**. The bit position can be specified directly by the **pos** field, or indirectly through the Shift Amount Register (SAR) if the **A** option is set.

If the **V** bit is set, the test succeeds when the selected bit is 1. If the **V** bit is clear, the test succeeds when the selected bit is 0.

When the condition is satisfied, execution branches to the target address, computed by sign-extending the offset, shifting it left by two bits, and adding it to the current instruction address. If the condition is not satisfied, execution continues with the next sequential instruction.

Operation

```
if ( instr.[A] ) pos <- SAR;
else           pos <- instr.[pos];

bVal <- testBit( RegR, pos );

cond <- ( instr.[V] ? (bVal == 1) : (bVal == 0) );

if ( cond ) PSW.[IA] <- addOfs32( PSW.[IA], ofs << 2 );
else       PSW.[IA] <- addOfs32( PSW.[IA], 4 );
```

Exceptions Taken Branch Trap
Instruction TLB Miss Trap

Notes None.

BE - Branch External

Syntax `BE [ ofs ] ( RegB ) [, RegR ]`

Format The BE instruction uses the signed immediate S15 format.

2	4	RegR	0	RegB	ofs
2	4	4	3	4	15

Description The BE instruction performs an unconditional branch to another segment. **RegB** contains a target segment and offset base to which the sign extended offset shifted by two is added. The return link is stored in **RegR**.

Because BE is a segment base relative branch, branching to a page with a different privilege level is possible. Branching from a lower privilege level to a higher one triggers an instruction protection trap. Branching from a higher privilege to a lower privilege level updates the privilege level in the processor status register. If no change is required, the privilege level remains unchanged.

Operation

<pre>RegR <- addOfs32(PSW.[IA], 4); PSW.[IA] <- addOfs32(RegB, RegX << 2);</pre>
--

Exceptions Taken Branch Trap
Instruction TLB Miss Trap

Notes None.

BR Branch Register

Syntax BR RegB [, RegR]

Format The BR instruction uses the immediate-15 instruction format.

2	3	RegR	0	RegB	0
2	4	4	3	4	15

Description The BR instruction performs an unconditional branch to a target address computed relative to the current instruction address segment (IA). The target address is formed by taking the contents of **RegB**, shifting it left by two bits, and adding it to the IA offset portion. The return link is stored in **RegR**.

Operation

```
RegR <- addOfs32( PSW.[IA], 4 );  
PSW.[IA] <- addOfs32( PSW.[IA], RegB << 2 );
```

Exceptions Taken Branch Trap
Instruction TLB Miss Trap

Notes None.

BV - Branch Vectored

Syntax BV RegB [, RegR, RegX]

Format The BV instruction uses the register instruction format.

2	4	RegR	1	RegB	0	RegX	0
2	4	4	3	4	2	4	9

Description The BV instruction performs an unconditional segment relative branch. The content of **RegX** is shifted by two and added to **RegB**, forming the target offset in the current segment. The return link is stored in **RegR**.

Because **BV** is a segment base relative branch, branching to a page with a different privilege level is possible. Branching from a lower privilege level to a higher one triggers an instruction protection trap. Branching from a higher privilege to a lower privilege level updates the privilege level in the processor status register. If no change is required, the privilege level remains unchanged.

Operation

```
RegR <- addOfs32( PSW.[IA], 4 );  
PSW.[IA] <- addOfs32( RegB, RegX << 2 );
```

Exceptions Taken Branch Trap
Instruction TLB Miss Trap

Notes None.

CBR - Compare and Branch

Syntax CBR.<cond> RegR, RegB, target

Format The CBR instruction uses the signed immediate S15 format:

2	9	Reg R	cond	Reg B	ofs
2	4	4	3	4	15

Description The CBR instruction compares the contents of **RegB** with **RegR** using the condition specified in the **cond** field. If the condition is satisfied, control is transferred to the instruction address computed by adding the sign-extended and word-aligned offset to the current instruction address. Otherwise, execution continues with the next sequential instruction. The comparison operation uses the same signed 64-bit comparator as the **CMP** instruction, and neither source register is modified.

Conditions The **cond** field encodes the branch condition as follows:

Value	Mnemonic	Branch if
0	EQ	RegR == RegB
1	LT	RegR < RegB
2	NE	RegR != RegB
3	GE	RegR >= RegB
4	LE	RegR <= RegB
5	reserved	
6	GT	RegR > RegB
7	reserved	

Operation

```
if ( cond( RegR, RegB ))
    psw.[IA] <- addOfs32( psw.[IA], ( ofs << 2 ));
else
    psw.[IA] <- addOfs32( psw.[IA] + 4 );
```

Exceptions Taken Branch Tap.
Instruction TLB Miss.

Notes The CBR instruction extends the **CMP** functionality by integrating a conditional branch. With a mode field of three bits available for the **cond** field, all six standard relational conditions (EQ, NE, LT, LE, GT, GE) are directly encoded, avoiding the need for operand swapping or result inversion.

See also the instruction notes on comparison and branch Condition encoding in the appendix.

CMP - Compare

Syntax

```
CMP RegR, RegB, RegA
CMP RegR, RegB, val

CMP[.D/W/H/B] RegR, ofs ( RegB )
CMP[.D/W/H/B] RegR, RegX ( RegB )
```

Format

The CMP instruction uses the register, the signed immediate S15, the indexed and the register indexed instruction formats.

0	6	RegR	cond	0	RegB	0	RegA	0
2	4	4	2	1	4	2	4	9

0	6	RegR	cond	1	RegB	val		
2	4	4	2	1	4	15		

1	6	RegR	cond	0	RegB	dw	ofs	
2	4	4	2	1	4	2	13	

1	6	RegR	cond	1	RegB	dw	RegX	0
2	4	4	2	1	4	2	4	9

Description

The **CMP** instruction compares two signed 64-bit operands. In register mode, it compares **RegB** and **RegA**; in immediate mode, it compares **RegB** and the immediate value; and in memory-reference formats, it compares **RegR** with the memory operand. The operands are called op1 and op2. The result of the comparison is stored in **RegR**. The instruction does not cause an arithmetic overflow but indexed instruction formats may cause a memory-reference-related trap.

Conditions

The **cond** field encodes the branch condition as follows:

Value	Mnemonic	Branch if
0	EQ	op1 == op2
1	LT	op1 < op2
2	NE	op1 != op2
3	GE	op1 >= op2

Operation

```
RegR <- compare( RegB, RegA );      (register format)
RegR <- compare( RegB, immOperand() ); (immediate format)
RegR <- compare( RegR, memOperand() ); (indexed formats)
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

The **CMP** instruction defines four comparison conditions (**EQ**, **NE**, **LT**, and **GE**). This choice minimizes hardware complexity, as each comparison can be evaluated directly from the ALU subtraction result without additional logic. The **LT** and **GE** operations are exact logical complements, and together with **EQ** and **NE**, form a complete and symmetric comparison set.

The remaining relational conditions (**LE** and **GT**) are not separate opcodes. In the register form of the instruction, they can be synthesized by the assembler through operand swapping ($A > B \iff B < A$, $A \leq B \iff B \geq A$). For addressing modes where operand swapping is not possible, these conditions can be generated by inverting the result of a corresponding **CMP** instruction.

See also the instruction notes on comparison and branch Condition encoding in the appendix.

DEP - Deposit Bit Field

Syntax

```
DEP[.Z/I] RegR, RegB, pos, len
DEP[.Z/I] RegR, RegB, SAR, len
```

Format The DEP instruction uses the special-13 instruction format.

0	7	RegR	1	RegB	I	A	Z	pos	len
2	4	4	3	4	1	1	1	6	6

Description The DEP instruction deposits a bit field from general register **RegB** into **RegR** at the specified position.

The rightmost bit of the target field is specified by the **pos** instruction field, or by the Shift Amount Register (**SAR**) if the **A** bit is set. The length of the field is specified by the **len** instruction field.

If the **Z** bit is set, the target register **RegR** is cleared before depositing the field. If the **I** bit is set, the value to deposit is taken from the **RegB** field rather than a register in **RegB**.

Operation

```
if ( instr.[Z] ) RegR <- 0;

if ( instr.[A] ) pos <- SAR;
else           pos <- instr.[pos];

if ( instr.[I] )
    RegR <- depositImm( RegR, RegB, pos, instr.[len] );
else
    RegR <- depositField( RegR, RegB, pos, instr.[len] );
```

Exceptions None

Notes None

DIAG - Diagnose Instruction

Syntax `DIAG id, RegB, RegA`

Format The DIAG instruction uses the signed immediate S9 instruction format.

3	15	RegR	id1	RegB	id2	RegA	0
2	4	4	3	4	2	4	9

Description The DIAG instruction provides a mechanism for diagnostic and system level operations.

The concatenated `id1` and `id2` fields specify the diagnostic function. The `RegB` and `RegA` fields may provide operands or data required by the diagnostic function. Any result is returned in `RegR`.

Operation

`RegR <- diagOperation( id, RegB, RegA );`

Exceptions Exceptions depend on diagnostic function and operands.

Notes None

DSR Double Shift Right

Syntax DSR RegR, RegB, RegA, amt
DSR RegR, RegB, RegA, SAR

Format The DSR instruction uses the special-9 instruction format.

0	7	RegR	2	RegB	0	A	RegA	0	amt
2	4	4	3	4	1	1	4	3	6

Description The **DSR** instruction concatenates the contents of registers **RegB** (left) and **RegA** (right) and performs a right shift operation. The lower 64 bits of the shifted result are stored in **RegR**.

The shift amount is specified by the **amt** instruction field, or by the Shift Amount Register (**SAR**) if the **A** bit is set.

Operation

```
if ( instr.[A] ) tmp <- SAR;
else             tmp <- instr.[amt];

RegR <- doubleShiftR( RegB, RegA, tmp );
```

Exceptions None

Notes None

EXTR - Extract Bit Field

Syntax

```
EXTR[.S] RegR, RegB, pos, len  
EXTR[.S] RegR, RegB, SAR, len
```

Format

The EXTR instruction uses the special-9 instruction format.

0	7	RegR	0	RegB	0	A	S	pos	len
2	4	4	3	4	1	1	1	6	6

Description

The EXTR instruction extracts a bit field from general register **RegB**.

The rightmost bit of the field is specified by the **pos** instruction field. If the **A** bit is set, the position value is taken from the Shift Amount Register (**SAR**) instead of the instruction field.

The length of the field is specified by the **len** instruction field. The extracted bit field is stored right-justified in **RegR**. If the **S** bit is set, the extracted value is sign-extended.

Operation

```
if ( instr.[A] ) tmp <- SAR;  
else           tmp <- instr.[pos];  
  
RegR <- extractField( RegB, tmp, instr.[len] );  
  
if ( instr.[S] ) RegR <- signExtend( RegR, instr.[len] );
```

Exceptions

None

Notes

None

FxCA - flush from data cache

Syntax FICA RegR, [RegX] (RegB)
 FDCA RegR, [RegX] (RegB)

Format The FICA instruction uses the standard indexed instruction format.

3	5	RegR	1	RegB	0
2	4	4	3	4	15

Description The FCA instruction flushes a cache line, selected by the effective address in RegB. This is a privileged instruction.

Operation RegR <- flushFromCache(RegB);

Exceptions Privilege violation trap.

Notes We may enhance this instruction with another register to for the target address by adding RegB with RegX, so that we can better support flushing cache ranges.

// ??? also uses the index and register indexed formats ?

IxTLB Insert TLB Entry

Syntax IITLB RegR, RegB, RegX
IDTLB RegR, RegB, RegX

Format The IITLB instruction uses the following instruction format.

3	4	RegR	0	RegB	0	RegX	0
2	4	4	3	4	2	4	9

The IDTLB instruction uses the following instruction format.

3	4	RegR	1	RegB	0	RegX	0
2	4	4	3	4	2	4	9

Description The ITLB instruction inserts a new entry into the translation lookaside buffer (TLB). The effective address is stored in **RegB** and additional control information is provided by **RegX**. This is a privileged instruction.

TLB Arguments **RegB** and **RegA** contain the arguments passed to the TLB subsystem. **RegB** contains the virtual address. The address is always o a page boundary and in the case of address ranges marks the starting page.

0	VPN	0
12	40	12

RegA contains the flags and the physical address of the address range.

Flags	0	Acc	Size	PPN	0
8	12	4	4	24	12

Operation **RegR** <- insertIntoTlb(**RegB**, **RegA**);

Exceptions Privilege Violation Trap.

Notes None.

LD Load Register

Syntax LD[.D/W/H/B/M] RegR, SignedImmed13(RegB)
LD[.D/W/H/B/M] RegR, RegX(RegB)

Format The LD instruction uses the indexed and register-indexed instruction formats.

1	7	RegR	0	M	0	RegB	dw	ofs	
2	4	4	1	1	1	4	2	13	

1	7	RegR	0	M	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description The LD instruction loads a data value from memory into a general-purpose register. Two formats are available: the immediate-offset format and the register-indexed format.

In the immediate-offset format, the signed immediate offset is shifted right by the **dw** field of the instruction and added to **RegB** to form the effective address. In the register-indexed format, **RegX** is added to **RegB** to form the effective address.

Optionally, the computed effective address may also be written back to the base register **RegB**. When this option is enabled, the sign of the offset determines the update behavior: A negative offset causes the base register itself to be used as the fetch address. A positive offset causes the sum of the base register and offset to be used. In both cases, the base register is updated.

The fetch address must be aligned to the data width specified. The value read from memory is sign-extended and placed in general register **RegR**.

Operation **RegR** <- memLoad(memOperand(instr));

Exceptions

- Data Alignment Trap
- Memory Access Violation Trap
- Memory Protection Violation Trap
- Data TLB Miss Trap

Notes The option to update the base register is often used in loops to fetch data from memory.

LDIL Load Immediate Left

Syntax LDIL[.L/.S/.U] RegR, val

Format The LDIL instruction uses the immediate-20 instruction format.

0	9	RegR	mode	val
2	4	4	2	20

Description The LDIL instruction deposits an immediate value into general register RegR. The optional suffix selects how the immediate field is positioned within the register. If no suffix is given, the default is .L mode (mode value = 1).

- .L (mode = 1) Deposits 20 bits of val into bits [31..12] of RegR.
- .S (mode = 2) Deposits 20 bits of val into bits [51..32] of RegR.
- .U (mode = 3) Deposits the lower 12 bits of val into bits [63..52] of RegR.

A mode value of zero is reserved and must not be used.

Operation

```
RegR <- ((val & 0xFFFF) << 12) // .L mode (1)
RegR <- ((val & 0xFFFF) << 32) // .S mode (2)
RegR <- ((val & 0xFFF) << 52) // .U mode (3)
```

Exceptions None.

Notes None.

LDO Load Offset

Syntax LDO [.B/H/W/D] RegR, ofs(RegB)

Format The LDO instruction uses the immediate-15 instruction format.

0	10	RegR	0	RegB	dw	ofs
2	4	4	3	4	2	13

Description The LDO instruction loads a value from memory using an offset addressing mode. The **ofs** value is sign extended and shifted by the **dw** field. The result is calculated by adding this value to the contents of **RegB**. The loaded value is then stored in **RegR**.

Operation

$\text{RegR} \leftarrow \text{addOfs32}(\text{RegB}, \text{signExtend}(\text{ofs}) \ll 2);$

Exceptions None.

Notes None

LDR Load Register Reserved

Syntax LDR RegR, SignedImmed13(RegB)

Format The LDR instruction uses the indexed instruction format.

1	9	RegR	0	RegB	3	ofs
2	4	4	3	4	2	13

Description The LDR instruction loads a value from memory into a general-purpose register and marks the address as "reserved" for a subsequent conditional store. The instruction is only the first step of an atomic read-modify-write pair. The subsequent STC (store conditional) instruction will succeed only if the reservation is still valid.

The signed immediate offset is shifted right by 3 and added to RegB to form the effective address.

Operation

```
RegR <- memLoad( memOperand( instr ));  
markAddressReserved(memOperand( instr ));
```

Exceptions

- Memory Access Violation Trap
- Memory Protection Violation Trap
- Data TLB Miss Trap

Notes This instruction is part of a "load-reserved / store-conditional" pair for atomic memory updates and often used for implementing locks or atomic counters. Reservation may be cleared by other stores to the same address by another processor.

LPA Load Physical Address

Syntax `LPA[.D/W/H/B] RegR, RegX( RegB )`

Format The LD instruction uses the indexed and register-indexed instruction formats.

3	2	RegR	0	M	1	RegB	0	RegX	0
2	4	4	1	1	1	4	2	4	9

Description The LPA (Load Physical Address) instruction computes the physical address corresponding to the given virtual address and loads it into **RegR**. The virtual address is formed by adding **RegX** to **RegB**. If there is no translation, a value of zero is returned.

The physical address, if found, is returned from information in the TLB. No memory access is performed. If there is a separate instruction and data TLB, the data TLB is used. For direct mapped segments the virtual address is the physical address.

Optionally, the computed effective address may also be written back to the base register **RegB**. When this option is enabled, the sign of the offset determines the update behavior: A negative offset causes the base register itself to be used as the fetch address. A positive offset causes the sum of the base register and offset to be used. In both cases, the base register is updated.

Operation `RegR <- translateAdr( addOfs32( RegB + RegX ) );`

Exceptions

- Alignment Trap
- Privileged instruction trap.
- Non access data TLB trap.

Notes None.

MBR - Move and Branch

Syntax MBR.<cond> RegR, RegB, target

Format The MBR instruction uses the following format:

2	10	RegR	cond	RegB	ofs
2	4	4	3	4	15

Description The MBR instruction copies the contents of **RegB** into **RegR**. If the condition specified in the **cond** field evaluates true for the new value in **RegR**, execution continues at the current instruction address plus the sign-extended offset. Otherwise, execution resumes with the next instruction.

Conditions The **cond** field encodes the branch condition as follows:

Value	Mnemonic	Branch if
0	EQ	RegR == 0
1	LT	RegR < 0
2	NE	RegR != 0
3	GE	RegR >= 0
4	LE	RegR <= 0
5	EV	RegR.[0] == 0
6	GT	RegR > 0
7	OD	RegR.[0] != 0

Operation

```
RegR <- RegB;
if ( cond( RegR )) PSW.[IA] <- addOfs32( PSW.[IA], signExtend(
    ofs ) << 2 );
else PSW.[IA] <- PSW.[IA] + 4;
```

Exceptions None.

Notes The CBR instruction extends the **CMP** functionality by integrating a conditional branch. With a mode field of three bits available for the **cond** field, all six standard relational conditions (EQ, NE, LT, LE, GT, GE) are directly encoded, avoiding the need for operand swapping or result inversion.

See also the instruction notes on comparison and branch Condition encoding in the appendix.

MFIA - Move from Instruction Address

Syntax MFIA [.U/S/L/A]RegR

Format The MFIA instruction uses the register instruction format.

3	1	RegR	1	mode	0
2	4	4	1	2	19

Description The MFIA instruction copies the current instruction address to register RegR. The mode field allows to select a portion of the instruction address.

- .A (mode = 0) copies IA to RegR. This is the default operation.
- .L (mode = 1) Deposits 20 bits of val into bits [31..12] of RegR.
- .S (mode = 2) Deposits 20 bits of val into bits [51..32] of RegR.
- .U (mode = 3) Deposits the lower 12 bits of val into bits [63..52] of RegR.

Operation

```
RegR <- PSW.[IA] // .A mode (0)
RegR <- ((PSW.[IA] & 0xFFFFF) << 12) // .L mode (1)
RegR <- ((PSW.[IA] & 0xFFFFF) << 32) // .S mode (2)
RegR <- ((PSW.[IA] & 0xFFF) << 52) // .U mode (3)
```

Exceptions None.

Notes It would be beneficial to have the possibility to just a subject of the instruction address. How about options to get the segment, the page number in the segment and the status bits. It is very similar to LDIL fields.

MFCR Move From Control Register

Syntax MFCR RegR, CReg

Format The MFCR instruction uses the control-register instruction format.

3	1	RegR	0	CReg	0
2	4	4	3	4	15

Description The MFCR instruction copies the contents of control register CRegB to general register RegR. Some control registers are privileged. Attempting to read them in user mode results in a privilege exception.

Operation

RegR <- CReg;

Exceptions - Privileged Operation Trap for some control registers.

Notes None

MTCR Move To Control Register

Syntax MTCR CReg, RegR

Format The MTCR instruction uses the control-register instruction format.

3	1	RegR	1	CReg	0
2	4	4	3	4	15

Description The MTCR instruction copies the contents of general register **RegR** to control register **CReg**. Some control registers are privileged. Attempting to write them in user mode results in a privilege exception.

Operation

CReg <- RegR;


Exceptions Privilege operation trap for some control registers.

Notes None

NOP No Operation

Syntax NOP

Format The NOP instruction uses no format.



Description The NOP (No Operation) instruction does nothing. It is typically used for timing, alignment, or to occupy a pipeline slot.

Operation noOperation();

Exceptions None.

Notes None.

OR Boolean Operation

Syntax

```
OR[.C/N] RegR, RegB, RegA
OR[.C/N] RegR, RegB, val

OR[.C/N/D/W/H/B] RegR, ofs( RegB )
OR[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The OR instruction uses the register, the immediate-19, the indexed and the register indexed instruction formats.

0	4	RegR	N	C	0	RegB	0	RegA	0
2	4	4	1	1	1	4	2	4	9

0	4	RegR	N	C	1	RegB	val		
2	4	4	1	1	1	4	15		

1	4	RegR	N	C	0	RegB	dw	ofs	
2	4	4	1	1	1	4	2	13	

1	4	RegR	N	C	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description

The **OR** instruction performs a bitwise OR operation on two 64-bit operands and stores the result in the target register **RegR**. The **AND** instruction performs a bitwise AND operation on two 64-bit operands and stores the result in the target register **RegR**. The second operand, **RegB** or the loaded content, can be negated by setting the "C" bit, and the result can be negated by setting the "N" bit. This instruction supports the immediate, register, and both memory-reference formats. The Boolean operation itself does not trigger any traps, but the indexed formats may cause a memory-reference related trap.

Operation

```
tmp <- RegA;                (register format)
tmp <- immOperand( instr );  (immediate format)
tmp <- memOperand( instr );  (indexed formats)

if ( instr.[C] ) tmp <- not( tmp );
RegR <- RegB | tmp;
if ( instr.[N] ) RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

PxCA - Purge from Cache

Syntax PICA RegR, RegB
 PDCA RegR, RegB

Format The PCA instruction uses the standard indexed instruction format.

3	5	RegR	0	M	0	RegB	mode	RegX	0
2	4	4	1	1	1	4	2	4	9

Description The PICA and PDCA instructions purge a cache line, selected by the effective address in RegB. This can be used by privileged code for cache management. Unlike FDCA, purge may also invalidate dirty lines without writing them back. The mode field indicates the cache type. A value of 0 refers to the instruction cache, a value of one refers to the data cache. A value of two refers to a unified cache.

Operation RegR <- purgeFromCache(addOfs32(RegB, RegX);

Exceptions None.

Notes Typically only allowed in supervisor mode.

// ??? also uses the index and register indexed formats ?

PRB Probe Virtual Address

Syntax

PRB RegR, RegB, RegA
PRB RegR, RegB, mode

Format

The PRB instruction uses both the register format and the special-9 format.

3	3	RegR	0	RegB	0	RegA	0
2	4	4	3	4	2	4	9

3	3	RegR	1	RegB	0	mode	0
2	4	4	3	4	4	2	9

Description

The PRB instruction probes the Translation Lookaside Buffer (TLB) to determine whether the effective address in **RegB** has a valid mapping. The result of the probe is returned in **RegR**: 0 indicates that no mapping exists, while 1 indicates a valid mapping.

In the register format, the access mode is specified in **RegA**. In the special-9 format, the mode is encoded directly in the instruction.

The mode specifies the type of access being probed:

Value	Access type
0	Read access
1	Write access
2	Execute access
3	Undefined

Operation

```
RegR <- probeVirtualAddress(RegB, RegA.[62..63]); // Reg  
RegR <- probeVirtualAddress(RegB, mode);           // Mode
```

Exceptions

- Non-access TLB Trap
- Privileged Operation Trap

Notes

Typically used by the operating system for virtual memory management.

PxTLB Purge TLB Entry

Syntax PITLB RegR, RegB, RegA
 PDTLB RegR, RegB, RegA

Format The PITLB instruction uses the indexed instruction format.

3	4	RegR	2	RegB	0	RegX	0
2	4	4	3	4	2	4	9

 The PDTLB instruction uses the indexed instruction format.

3	4	RegR	3	RegB	0	RegX	0
2	4	4	3	4	2	4	9

Description The PITLB and PDTLB instructions purge (removes) an entry from the translation lookaside buffer (TLB). The entry to remove is identified by the effective address formed from **RegB** and **RegX**. This is privileged instruction.

Operation RegR <- purgeFromTlb(RegB, RegA);

Exceptions - Privileged operation trap.

Notes None.

// ??? standard memory address formats ?

RFI Return from Interrupt

Syntax RFI [RegR]

Format The RFI instruction uses the branch instruction format.

3	7	RegR	0	0
2	4	4	3	19

Description The RFI instruction restores processor state after an interrupt or exception handler has completed. Execution resumes at the program counter saved during the interrupt and the processor state (including privilege level, interrupt enable flags, and other PSR fields) is restored from the saved copy.

Operation PSW <- IPSW;

Exceptions Privileged operation trap.

Notes None.

// ??? what is RegR used for ?

RSM Reset System Mask

Syntax RSM RegR

Format The RSM instruction uses the system instruction format.

3	6	RegR	0	0	0	0	0	0	0
2	4	4	3	13	1	1	1	1	1

Description The RSM (Reset System Mask) instruction clears bits in the system mask register that are set in RegR.

Operation `SystemMask <- SystemMask & instr.[5:0];`

Exceptions - Privileged instruction trap.

Notes None.

// ?? what is RegR used for ? return old mask ?

SHLxA Shift Left and Add

Syntax

SHLxA RegR, RegB, RegA
SHLxA RegR, RegB, SignedImmed13

Format

The SHLxA instruction supports both register and immediate formats.

0	8	RegR	amt	0	RegB	0	RegA	0
2	4	4	2	1	4	2	4	9

0	8	RegR	amt	0	RegB	2	val
2	4	4	2	1	4	2	13

Description

The SHLxA instruction shifts the contents of **RegB** left by the amount specified in the **amt** field. In register format, **RegA** is added to the shifted value. In immediate format, if the I bit is set, the sign-extended **val** field is added instead. The resulting value is stored in **RegR**.

Operation

RegR <- (RegB << instr.[amt]) + RegA; (Register)
RegR <- (RegB << instr.[amt]) + immOperand(instr);(Immediate)

Exceptions

Integer Overflow Trap

Notes

None

SHRxA - Shift Right and Add

Syntax

```
SHRxA RegR, RegB, RegA
SHRxA RegR, RegB, val
```

Format The SHRxA instruction supports both register and immediate formats.

0	8	RegR	amt	1	RegB	0	RegA	0
2	4	4	2	1	4	2	4	9

0	8	RegR	amt	1	RegB	2	val
2	4	4	2	1	4	2	13

Description The SHRxA instruction shifts the contents of **RegB** right by the amount specified in the **amt** field. In register format, **RegA** is added to the shifted value. In immediate format, if the I bit is set, the sign-extended **val** field is added instead. The resulting value is stored in **RegR**.

Operation

$\text{RegR} \leftarrow (\text{RegB} \gg \text{instr.}[\text{amt}]) + \text{RegA}; \quad (\text{Register})$ $\text{RegR} \leftarrow (\text{RegB} \gg \text{instr.}[\text{amt}]) + \text{immOperand}(\text{instr}); (\text{Immediate})$
--

Exceptions Integer Overflow Trap

Notes None

SSM Set System Mask

Syntax SSM RegR, mask

Format The SSM instruction uses the system instruction format.

3	6	RegR	1	0	0	0	0	0	0
2	4	4	3	13	1	1	1	1	1

Description The SSM (Set System Mask) instruction updates the system mask register. The previous mask is returned in RegR.

Operation `SystemMask <- SystemMask | instr.[5:0];`

Exceptions - Privileged instruction trap.

Notes None.

ST Store Register

Syntax `ST[.D/W/H/B] RegR, SignedImmed13( RegB )`
 `ST[.D/W/H/B] RegR, RegX( RegB )`

Format The ST instruction uses the indexed and register-indexed instruction formats.

1	8	RegR	0	M	0	RegB	dw	S-IMM-13	
2	4	4	1	1	1	4	2	13	

1	8	RegR	0	M	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description The ST instruction stores a value from a general-purpose register into memory.

In the immediate-offset format, the signed immediate offset is shifted right by the **dw** field of the instruction and added to **RegB** to form the effective address. In the register-indexed format, **RegX** is added to **RegB** to form the effective address.

Optionally, the computed effective address may also be written back to the base register **RegB**. When this option is enabled, the sign of the offset determines the update behavior: A negative offset causes the base register itself to be used as the fetch address. A positive offset causes the sum of the base register and offset to be used. In both cases, the base register is updated.

The store address must be aligned to the data width specified. The value stored to memory is stored as is in the data width specified.

Operation `memStore( memOperand(), RegR );` (for indexed formats)

Exceptions - Data Alignment Trap
 - Memory Access Violation Trap
 - Memory Protection Violation Trap
 - Data TLB Miss Trap

Notes The jury is still out if the M bit is not too expensive in hardware. We would need to read three registers. One for the value, and up to two for the index calculation.

STC Store Register Conditional

Syntax STC RegR, SignedImmed13(RegB)

Format The STC instruction uses the indexed instruction format.

1	10	RegR	0	RegB	3	S-IMM-13
2	4	4	3	4	2	13

Description The STC instruction stores a value from a general-purpose register to memory only if the memory address is still reserved by a prior LDR instruction. - If the reservation is valid, the store succeeds and clears the reservation. - If the reservation has been lost (e.g., another processor wrote to the address), the store fails and the memory is unchanged.

Operation

```
if (isAddressReserved( memOperand( instr ))) {  
    memStore( memOperand( instr ), RegR );  
    clearReservation( memOperand( instr ));  
} else {  
    storeFailedFlag |= 1;  
}
```

Exceptions

- Data Alignment Trap
- Memory Access Violation Trap
- Memory Protection Violation Trap
- Data TLB Miss Trap

Notes This instruction is the second part of a "load-reserved / store-conditional" pair for atomic memory updates and commonly used for implementing spinlocks, mutexes, and atomic counters in multiprocessor systems.

SUB Integer Subtraction

Syntax

```
SUB RegR, RegB, RegA
SUB RegR, RegB, val

SUB[.D/W/H/B] RegR, ofs ( RegB )
SUB[.D/W/H/B] RegR, RegX ( RegB )
```

Format

The SUB instruction uses the register, the immediate-15, the indexed and the register indexed instruction formats.

0	2	RegR	0	RegB	0	RegA	0
2	4	4	3	4	2	4	9

0	2	RegR	1	RegB	val		
2	4	4	3	4	15		

1	2	RegR	0	RegB	dw	ofs	
2	4	4	3	4	2	13	

1	2	RegR	1	RegB	dw	RegX	0
2	4	4	3	4	2	4	9

Description

The SUB instruction subtracts two signed 64-bit operands and stores the result in the target register **RegR**. This instruction supports the immediate, register, and both memory-reference formats. An arithmetic overflow triggers a numeric overflow trap, and indexed instruction formats may cause a memory-reference-related trap.

Operation

```
RegR <- RegB - RegA;           ( register format )
RegR <- RegB - immOperand( instr ); (immediate format)
RegR <- RegR - memOperand( instr ); (indexed formats)
```

Exceptions

Integer Overflow Trap
Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

TRAP Trap Instruction

Syntax TRAP tid, regB, regA

Format The TRAP instruction uses the standard indexed instruction format.

3	14	0	tid1	RegB	tid2	RegA	0
2	4	4	3	4	2	4	9

Description The TRAP instruction generates a synchronous trap to handle exceptional conditions. The `tid1` and `tid2` are concatenated to form the trap Id. The instruction may be used to invoke operating system routines or exception handlers.

Operation

```
raiseTrap( id, RegB, RegA );
```

Exceptions The trap raised.

Notes None

XOR Boolean Operation

Syntax

```
XOR[.C/N] RegR, RegB, RegA
XOR[.C/N] RegR, RegB, val

XOR[.C/N/D/W/H/B] RegR, ofs( RegB )
XOR[.C/N/D/W/H/B] RegR, RegX ( RegB )
```

Format

The XOR instruction uses the register, the immediate-19, the indexed and the register indexed instruction formats.

0	5	RegR	N	0	0	RegB	0	RegA	0
2	4	4	1	1	1	4	2	4	9

0	5	RegR	N	0	1	RegB	val		
2	4	4	1	1	1	4	15		

1	5	RegR	N	0	0	RegB	dw	ofs	
2	4	4	1	1	1	4	2	13	

1	5	RegR	N	0	1	RegB	dw	RegX	0
2	4	4	1	1	1	4	2	4	9

Description

The XOR instruction performs a bitwise XOR operation on two 64-bit operands and stores the result in the target register **RegR**. The result can be negated by setting the "N" bit. This instruction supports the immediate, register, and both memory-reference formats. The Boolean operation itself does not trigger any traps, but the indexed formats may cause a memory-referencerelated trap.

Operation

```
RegR <- RegB ^ RegA;           (register format)
RegR <- RegB ^ immOperand();    (immediate format)
RegR <- RegR ^ memOperand();    (indexed formats)
if ( instr.[N] ) RegR <- not( RegR );
```

Exceptions

Data Alignment Trap
Memory Access Violation Trap
Memory Protection Violation Trap
Data TLB Miss Trap

Notes

None.

Synthetic Instructions

The Twin-64 instruction set provides a rich set of options for individual instructions. Setting a defined option bit in an instruction can enable additional functionality with minimal impact on the overall data path. Similarly, instructions that use general register zero as an argument can take advantage of predefined behaviors.

For better readability and convenience, an assembler can offer *synthetic instructions*, which are higher-level mnemonics generated from the base instruction set. These synthetic instructions simplify programming by providing a more natural or concise representation of commonly used instruction patterns.

Synthetic instructions are also useful when a single instruction cannot fully express the desired operation. For example, loading an immediate value larger than the instruction field allows requires an instruction sequence. An assembler can provide a generic "load immediate" synthetic instruction that automatically emits such an instruction sequence depending on the size of the value. Similarly, if the offset to a base register is too small to reach a desired data item, the assembler can transparently emit an additional 'ADDIL' instruction to compute the correct address.

The synthetic instruction set can be organized into several groups:

- **Immediate Operations:** yyy.
- **Register Operations:** yyy.
- **Arithmetic Operations:** yyy.
- **Shift and Rotate Operations:** yyy.
- **Memory Reference Operations:** yyy.

Immediate Operations

LDI

Register Operations

CLR

CPY

COM

NEG

Arithmetic Operations

INC

DEC

EXT8

EXT16

EXT32

Shift and Rotate Operations

ASR

LSR

ROR

ROL

Memory Reference Operations

LOAD

STOR

Part IV

I/O Subsystem

Input/Output

Concepts

Twin-64 features memory a mapped I/O architecture.

memory mapped

dedicated address space in physical memory

bus concept

module concept

modules have registers - load / store, however registers do not have to be 64 bit wide.

the system bus features up to 64 modules. that is: 64 HPA pages, need: 1Mb address range.

a module listens to three address ranges

- hard physical address range. a page. privileged.

- soft physical address range. allocated by the module for further space. aligned on a page boundary.

- broadcast address range. Mirrors the system bus range. A write to a register in that range applies to all corresponding registers of a module on the same bus.

modules are processors, memory and I/O cards.

we could keep processor stats in the HPA space

memory module has entire memory as SPA space

I/O Address Space

As shown in the overall physical memory layout, the physical memory address range dedicated to I/O modules is located from 0x00_F000_0000 to 0x00_FFFF_FFFF. The following figure depicts the I/O memory address space ranges in more details




Figure 15: I/O Memory Address Range

The I/O address range is a 256Mbyte area in the first 4 Gbyte segment, i.e. segment 0. The area is divided into several ranges. The first range contains the processor dependent code. There is a set of library functions to manage the processor itself but also provide low-level primitives for the boot process and the operating system.

Hard physical address range

The **hard physical address** range contains a I/O page for each I/O modules installed. There room for 64 I/O modules.

Broadcast address range

The **broadcast address** range mirrors the hard physical address range. However, a write to these locations is a write operation to which all I/O modules will react.

Soft physical address range

The **soft physical address** range is an area where space is allocated to an I/O module. The I/O module reacts to a memory access operation in its configured soft physical address range.

I/O Elements

I/O elements are 128 bytes, i.e. 16 regs

I/O elements can also be two pages, one priv and one user level, mapped...

Part V

Runtime Environment

Runtime Overview

No CPU architecture with an idea of a runtime environment. Although the instruction set can be used to implement many models of runtime environments, there are common principles that exploit the CPU architecture. In fact, several CPU concepts were selected with a runtime already in mind. This chapter will present the Twin-64 runtime environment and describe the high level system model, the register usage convention, the execution model and calling convention.

System Environment

start high level what the system is

a system runs tasks grouped into jobs

The following figure depicts a high level overview of a software system for Twin-64. At the center is the execution thread, which is called a **task**. A task will consist of the current execution object and the task data area. Tasks belong to a **job**. All tasks of a job share a common job data area. At the highest level is the **system**, which contains the global data area for the system.




Figure 16: Runtime Overview

Runtime Environment

basic concepts

a runtime at any time is a code and a data region.

a program or library, once prepared, is a read only / execute object.

the loader will update the LTAB for a program or library loaded.

once loaded, it will not be released during the life time of the system ? To be investigated.
Although we have a million regions, can we unload ? And if so, how do we make sure that the client LTAB data is updated properly ?

Register Usage

The CPU features general registers and control registers. As described before, all general registers can do computation as well as address computation.

global data - pointed to by DP (runtime architected to be a specific register ?)

stack data - pointed to by SP (runtime architected to be a specific register ?)

caller saved -

callee saved -

R0, R1 special

Task Execution

segment - code, data

global data - pointed to by DP

stack data - pointed to by SP

modules and functions

modules have global data

functions allocate a stack frame

stack grows upward

SP minus addressing

frame marker keeps previous SP and RL

function call - RP optionally stored in frame marker

first 4 args are passed in registers

remaining args in stack frame

Stack Layout

Note: sketches and notes for now...

Stack Frame




Figure 17: Stack Frame

Stack Frame Marker

stack frame marker has four double words. (make it a picture ...)




Figure 18: Stack Frame Marker


A stack frame marker could also be smaller, if we consider to put DP and the reserved word as local variables in the stack frame.

Calling Convention

Note: sketches and notes for now...

Local Call

By far the most used call...



Linkage Table

each module as a set of entries in the linkage table for the external references of this module.

The assembler / compiler generates the linkage table skeleton.

an external procedure has an entry, also a procedure label




Figure 19: Processor Configuration Register

flags need to be defined, important that the BE instruction always ignores the upper 12 bits.

table is at an architected location on the current region, offset zero.

the offset in the BE instruction refers to the entry. This allows for 16384 entries considering a one double word entry.

offset is the index into the linkage table

External Call

The external address is taken from the linkage table entry. The offset was computed at compile time.

```

1      .
2      .
3      .
4      .

```

An exported procedure will have a different entry and ext sequence compared to a local procedure. The `.ENTER` and `.LEAVE` pseudo instructions will generate the respective code sequences.

The `.CALLINFO` pseudo instruction will describe the procedure.

The code snippet below is the maximum code to be generated.

```

1      .CALLINFO ...      ; procedure description
2
3      ;ENTER sequence
4      ST SP,(SP)          ; save PSP
5      ST RL, -8(SP)       ; save RL
6      ST DP, -16(SP)      ; save DP
7      MFIA.S R1           ; get module LTAB address
8      LD DP, <ofs>(R1)     ; load module DP
9      LDO SP,<size>(SP)    ; allocate stack frame
10     .
11     .
12     .
13     .
14     ;LEAVE sequence
15     LDO SP, -<size>(SP) ; deallocate stack frame
16     LD DP, -16(SP)      ; load previous DP
17     LD RL, -8(SP)       ; load return link
18     BE (RL)             ; external branch return

```

the DP value is taken from the first entry in the LTAB section for the module.

for privilege level changes, there needs to be a stub on a gateway page.

```

1      ; on a gateway page
2      .
3      .
4      B local target ;jump to the target procedure.
5      .
6      .

```

to investigate: is a simple branch enough ? should it rather be a jump based on an entry in an export table ? LTAB alike ?

Part VI

Processor Design

Part VII

Simulator Environment

Part VIII

Appendix

Appendix - Instruction Commentary

Comparison and Branch Condition Encoding

This appendix explains the condition-code encoding and the implementation philosophy used by the **CMP**, **CBR**, **ABR** and **MBR** instructions. The design goals are to minimize hardware complexity by reusing a single subtraction/comparator path. Comparison operators are encoded in a compact and symmetric way that makes condition inversion trivial. The following instructions use the comparison operations.

- **CMP** The compare instruction uses a 2-bit **cond** field (four primitives) and produces a boolean result (0/1) in a register.
- **CBR** The compare-and-branch uses a 3-bit **cond** field (eight encodings) and branches based on comparison between two registers.
- **ABR** The add-and-branch performs an add operation and branches using a 3-bit **cond** field that tests the add result against zero or other simple predicates.
- **MBR** move-and-branch also uses the same 3-bit **cond** field and tests the loaded result.

All three branch instructions (**CBR**, **ABR**, **MBR**) share the same 3-bit condition-code encoding so decoder logic is uniform across branch forms. **CMP** uses a compact 2-bit encoding (four primitives) for hardware efficiency; the assembler and/or an **inv** operation are used to provide the remaining relations where needed.

Branch target computation uses a sign-extended 15-bit offset shifted left by two, allowing word-aligned branch destinations within a 64 KB range relative to the current instruction address.

Design Rationale Implementing just four compare primitives minimizes combinational logic in the arithmetic/comparator unit. The primitives chosen are **EQ**, **NE**, **LT** and **GE**. These four are symmetric in pairs: **EQ** / **NE** and **LT** / **GE**. Each pair is the logical complement of the other, so an *invert* operation (flip of a single decode bit) yields the complementary condition. That property makes both ISA decoding and assembler inversion trivial.

Where a comparison such as **GT** or **LE** is required, the assembler can safely swap operands (register form, no side-effects), it will rewrite the comparison using operand swap (e.g. **A > B** **B < A**). Where operand swapping is not possible (immediate or memory addressing modes), the condition is produced by inverting the result of an available primitive (either by an **inv** instruction that flips a 0/1 boolean, or by a branch invert bit in the branch encoding).

Encodings The compare instruction (2-bit **cond**) use the following compare modes:

Value	Mnemonic	Meaning
0	EQ	op1 == op2
1	LT	op1 < op2
2	NE	op1 != op2
3	GE	op1 >= op2

The 2-bit encoding is chosen so that the left most bit is the *invert* bit: flipping it yields the complementary condition for every pair.

The other branch type instructions use the following compare modes:

Value	Mnemonic	Branch if
0	EQ	Reg == 0
1	LT	Reg < 0
2	NE	Reg != 0
3	GE	Reg >= 0
4	LE	Reg <= 0
5	EV	Reg.[0] == 0
6	GT	Reg > 0
7	OD	Reg.[0] != 0

The 3-bit encoding is chosen so that middle bit is the *invert* bit. Flipping it yields the complementary condition for every pair. This symmetry keeps branch inversion extremely cheap in hardware and consistent across all branch forms.

Logical decoding and hardware signals All conditions used by the comparator/branch unit are derived from a few simple signals computed by the ALU/comparator after a subtraction (or directly from the result when testing against zero). The following Verilog code sequence depicts a possible implementation.

```

1  wire Z  = (result == 64'd0);
2  wire S  = result[63];
3  wire B0 = result[0];
4
5  wire cond_EQ = Z;
6  wire cond_LT = S;
7  wire cond_NE = ~Z;
8  wire cond_GE = ~S;
9  wire cond_LE = S | Z;
10 wire cond_GT = (~S) & (~Z);
11 wire cond_EV = ~B0;
12 wire cond_OD = B0;

```

The condition select is a multiplexer that select on of the wires above. Because all branch forms share the same condition encoding scheme, the branch unit uses a single small logic table to produce **branch_taken**, simplifying pipeline wiring and verification.

CMP Instruction Assembler rules and rewrites If the two operands are registers it is safe to use operand swapping in order to avoid generating extra instructions (invert or boolean complement) and to keep assembly readable. When a programmer writes a relation that is not one of the four primitives, the assembler rewrites it to:

`CMP.GT rR, rA, rB -> CMP.LT rR, rB, rA`

`CMP.LE rR, rA, rB -> CMP.GE rR, rB, rA`

Swapping is only applied when both operands are pure (register or immediate) and no addressing side-effects are present. If swapping is not safe (memory-reference, indexed forms, or side-effecting addressing modes), the assembler emits the corresponding primitive followed by an `invert` operation.

The branch type instructions encode all common relations directly, so assembler rewrites are not required.

Appendix - Instruction Notation Routines

describe the routines used in the instruction set operation part...

Table 6: Instruction Notation Routines

Function	Description
<code>addOfs32(base, ofs)</code>	...
...	...

Appendix - Programming Notes

xxx

Appendix -TLB and Caches

??? perhaps add to processor appendix...

Appendix -Diagnostic Functions

xxx

Table 7: Diagnostic Functions

Id	Function	Arg0	Arg1
0	No operation	0	0
1	Write status display	value	-

Appendix - Glossary

Address Space A logical domain of addresses. The CPU may define multiple address spaces, such as *main memory space* and *I/O space*. Each space defines an independent address range and access semantics.

Region A contiguous range of addresses within an address space, reserved for a particular purpose. Examples: *code region*, *data region*, *stack region*.

Object A unit of storage within a region. Objects correspond to program-visible entities or runtime allocations, such as global variables, heap blocks, or stack frames. Examples: *global object*, *heap object*, *stack object*.

ELF Section A named subdivision in an ELF file used by the linker and compiler (e.g., *.text*, *.data*, *.bss*). Sections are grouped by the linker into ELF segments.

ELF Segment A loadable portion of an ELF file (defined by a program header). Each ELF segment is mapped into one or more regions of an address space during program loading.

Mapped Region A runtime mapping that associates an ELF segment or memory source with a region in an address space. The emulator or loader tracks mapped regions to manage memory access and protection.

